

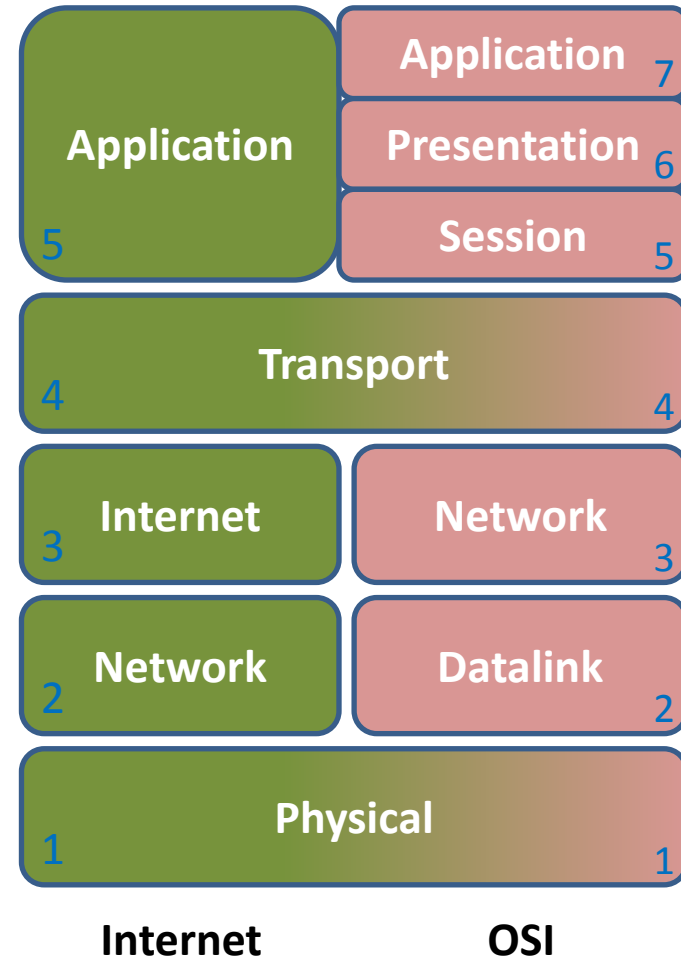
Internetworking

G54ACC – IP and Up

Lecture 1

Recap

- IP sits at layer 3
 - OSI *network layer*
 - TCP/IP *Internet layer*
- Approximates a *fully-connected network* abstraction
- IP packets are *hop-by-hop* routed by *store-and-forward* devices



Contents

- Overview
- Addressing
- Routing vs. Forwarding
- Address Management
- Control
- Evolution
- Alternatives & Summary

Contents

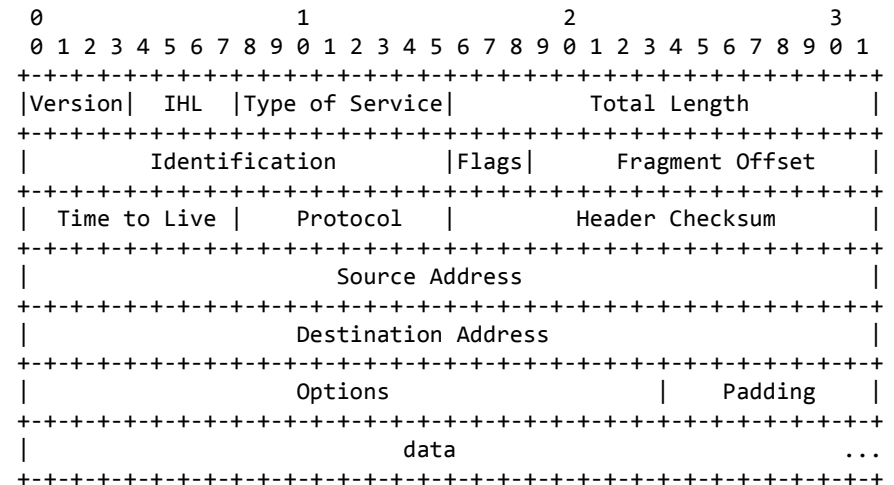
- Overview
 - Encapsulation
- Addressing
- Routing vs. Forwarding
- Address Management
- Control
- Evolution
- Alternatives & Summary

Overview

- Hosts are allocated addresses
- Generated data is encapsulated
- Packets passed to datalink for transmission
 - Commonly Ethernet in some form
 - *Historic interest: Token-Ring, ATM*
- Routers receive packets, examine, forward
 - May have non-IP devices between routers
 - Details of datalink are abstracted

Encapsulation

- Data comes down from above in chunks
 - Sockets API imposes buffers even on streams
 - These chunks may be too large for the datalink
 - Packetization generates suitable sized chunks
- IP header is prepended



RFC 791, Internet Protocol

<http://www.rfc-editor.org/rfc/rfc791.txt>

IEN 46, Addressing and Naming

<http://www.postel.org/ien/txt/ien46.txt>

Contents

- Overview
- Addressing
 - Address Aggregation
 - Classless Inter-Domain Routing
- Routing vs. Forwarding
- Address Management
- Control
- Evolution
- Alternatives & Summary

Addressing

- Host addresses in IPv4 are 32 bits long
 - Intended to be globally unique
 - Really it's *interfaces* that get allocated addresses
 - 32 bits is no longer enough
 - E.g., virtualization, mobile telephony
- Commonly depicted in 4 parts
 - 32 bits == 4 bytes
 - Dotted quad notation (msb first): a.b.c.d
 - E.g., 128.243.35.39 == 0x80f32327 == 2163417895

Address Aggregation

- Need to determine the *route* from the *address*
 - How do we get where we want to go?
 - Group addresses for efficiency: 2^{32} bits not enough but still a lot!
- Divide address into *network* and *host* parts
 - Hierarchical addressing
- Originally fixed: first 8 bits for network, remaining for hosts
 - What to do when >254 networks?
- Next, moved to *classful*: A, B, C, (D/multicast, E/reserved)
 - Indicate length of network part by high-order bits
 - But many wasted addresses as Internet grew

Classless Inter-Domain Routing

- Subnetting: make *netmask* explicit
 - Cluster addresses internally
- CIDR, RFC 1338
 - Gives rise to *prefixes* with explicit *prefix lengths*
- E.g., $172.16/12 == (172.16.0.0, 255.240.0.0)$
 - $A = x/8$, $B = x/16$, $C = x/24$ but required routing protocol changes

Contents

- Overview
- Addressing
- **Routing vs. Forwarding**
 - Forwarding
 - Other Forwarding Actions
- Address Management
- Control
- Evolution
- Alternatives & Summary

Routing vs. Forwarding

- Router receives an IP packet: what to do?
 - Drop or forward?
 - If forward, which port to transmit it on?
- Making this decision is **forwarding**
 - Which port gets it closer to the destination address?
 - IP bases this decision (almost) solely on `dstip`
- Building up the information to do so is **routing**
 - Where are all the addresses at the moment?

Forwarding

- Need to map packet's destination to port
- Routing generates (prefix, port) mapping
 - The *forwarding table*
- Map *address* to *prefix* via *longest prefix match*
 - Means the *most specific* entry is used
 - If no match, then use *default route*; else drop
- How to structure forwarding table?
 - Minimise lookup latency
 - Handle variable and substantial update dynamics

Other Forwarding Actions

- Decrement the *time-to-live* (TTL) field
 - Ensures that looping packets eventually die
- Perform any *fragmentation* necessary
 - IP abstracts datalinks with different max PDUs
 - Generate N packets from the first, setting fragment offset
 - ...unless Don't Fragment is set: Path MTU Discovery

Contents

- Overview
- Addressing
- Routing vs. Forwarding
- **Address Management**
 - Macro – Internet-wide
 - Micro – Network-local
- Control
- Evolution
- Alternatives & Summary

Address Management: Macro

- Internet Assigned Numbers Authority (IANA)
 - Co-ordinates number spaces
- Delegates *netblocks* to Regional Internet Registries (RIRs)
 - Africa, Asia/Pacific, North America, Latin America, EMEA
 - ...and down to National and Local registries (NIRs, LIRs)
- Approach appropriate registry for an allocation
 - Must provide suitable justification
 - Much harder to get a large allocation (== short prefix)
- Deeply manual process involving much politics
 - <http://www.iana.org/numbers/> and <http://www.iana.org/>

Address Management: Micro

- Used to also be very manual
 - A big text file containing (IP, Ethernet) address map
 - Also needed to maintain subnet allocations for routing protocols
- Dynamic Host Configuration Protocol, DHCP
 - RFC 2131
 - “Can anyone give me an address?”
 - Usually provides a pile of other configuration information
- Address Resolution Protocol, ARP
 - “Does anyone know who’s at <ipaddr>?”
- Both utilise IP broadcast address, 255.255.255.255

Contents

- Overview
- Addressing
- Routing vs. Forwarding
- Address Management
- **Control: ICMP**
- Evolution
- Alternatives & Summary

Internet Control Message Protocol

- ICMP: RFC 792
 - In-band control: uses IP messages
- Most common operations:
 - ECHO REQUEST/RESPONSE (ping – “are you there?”)
 - UNREACHABLE (host, network, gateway)
 - TIME EXCEEDED (TTL was decremented to zero)
- Implementation of *traceroute*
 - How to determine route your packets take?
 - Transmit packet to destination with increasing TTL
 - Receive ICMP time exceeded from each enabled gateway

Contents

- Overview
- Addressing
- Routing vs. Forwarding
- Address Management
- Control
- **Evolution**
 - Multicast
 - Mobile
 - IPv6
- Alternatives & Summary

Multicast

- Have considered *unicast* (1:1) and *broadcast* (1:all)
- IP natively supports *multicast* (M:N)
 - Efficiency gains if the routers do the replication
 - Note multicast support exists at some lower layers too...
- Everything becomes more complex
 - Need to manage *group membership*
 - Need to manage routing between groups
- For example,
 - IGMP controls local group membership
 - PIM, MOSPF handle routing between local groups
 - MBGP permits different routing between networks

Mobile

- IP can natively support mobility; RFC3344
 - Again, so do some lower layers...
- Assign *Care-of-Address* in addition to home address
 - *Home-agent* and *foreign-agent* manage nodes
- Again, adds complexity: protocols for
 - Authenticated registration to inform HA of CoA
 - Extension to ICMP for agent discovery
 - Routing protocols, typically tunnelling
- What about multicast? (ARP? Broadcast? ...)

IPv6

- IPv4 is dead! Long live IPv6!
 - Well, sort of: in transition since 1998
- Primary benefit: larger address space (128b)
 - But what about NAT
- Other benefits:
 - Integrated security, back-ported as IPSec
 - Various auto-configuration mechanisms
 - Mobility support: avoid *triangular routing*
 - Jumbograms, extended options space, ...
- Costs: all the other protocols need to change too!
 - *Network layer is **much more** than just an encapsulation*

Contents

- Overview
- Addressing
- Routing vs. Forwarding
- Address Management
- Control
- Evolution
- **Alternatives & Summary**

Alternatives

- Packet systems: CLNP, CLNS
 - Actually still used, a little! IS-IS/CLNP/802.3
- Circuit switching
 - A real actual copper circuit (old telephones)
 - Digital telephony implies framing (SONET, &c.)
 - Starting to move onto IP network anyway
- Statistical multiplexing with circuits
 - Virtual circuits (ATM)
 - Flow switching (Ipsilon)
 - Compare circuit setup vs. Per-packet routing

Summary

- IP provides encapsulation and packetization of data from higher layers
- Packets are forwarded through the network toward the specified destination by routers
- Building up information to forward is routing
 - *Topic of next lecture*
- Several control protocols, processes and organizations provide IP infrastructure