

Routing

G54ACC – IP and Up

Lecture 2

Recap

- “The Internet” consists of connected routers
- Routers *store’n’forward* packets toward destinations
- Decisions are based on IP destination address and contents of *routing tables*

Contents

- Overview
- Link-state routing
- Distance-vector routing
- Inter-network routing
- Alternatives & Summary

Contents

- Overview
 - What is routing?
 - IP routing
- Link-state routing
- Distance-vector routing
- Inter-network routing
- Alternatives & Summary

What is Routing?

- The process of building up information to enable forwarding
 - How does a router figure out the correct port on which to forward a packet?
 - Implicit: “correct” means “most efficient”
 - ...subject to other constraints
- Why is it a problem?
 - Scalability: networks may become large
 - Dynamics: need to handle host and link failures

IP Routing

- Two basic techniques for wireline IP:
 - *Link-state* routing
 - *vs. Distance-vector* routing
- Both equivalent but make different tradeoffs
- Both require some degree of coordination
- Many other techniques in general
 - Particularly in ad-hoc wireless and other networks
 - Geographical and map-based are common
 - What information is reliably available?

Ok, *three* basic techniques

- Static routing
 - Entries in routing table independent of network state
 - Entered manually, or via DHCP
 - Common in hosts and very very small networks
- For example:

```
[Linux 2.6] $ route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
10.8.35.0	*	255.255.255.0	U	0	0	0	eth2
192.168.9.0	*	255.255.255.0	U	0	0	0	br0
default	10.8.35.1	0.0.0.0	UG	100	0	0	eth2

A More Complex Example

```
[Mac OSX] $ netstat -rlf inet
```

Routing tables

Internet:

Destination	Gateway	Flags	Refs	Use	Mtu	Netif	Expire
default	tfa1-gw-v-35-35-0.	UGSc	38	0	1500	en0	
default	link#7	UCSI	0	0	1500	en3	
127	localhost	UCS	0	0	16384	lo0	
localhost	localhost	UH	5	56699	16384	lo0	
128.243.35/24	link#4	UCS	5	0	1500	en0	
tfa1-gw-v-35-35-0.	0:18:74:1e:bd:40	UHLWI	38	13	1500	en0	345
ppshorizon316.nott	0:25:64:9c:d9:61	UHLWI	0	665	1500	en0	1183
xpshorizoncanon.no	0:1e:8f:2e:de:8f	UHLWI	0	0	1500	en0	1022
puidhcp-035-215.is	0:23:18:c0:67:7e	UHLWI	0	0	1500	en0	1196
puidhcp-035-223.is	localhost	UHS	0	0	16384	lo0	
128.243.35.255	ff:ff:ff:ff:ff:ff	UHLWbI	0	8	1500	en0	
169.254	link#4	UCS	1	0	1500	en0	
greyjay.local	localhost	UHS	0	0	16384	lo0	
169.254.255.255	link#4	UHLW	1	113	1500	en0	

Contents

- Overview
- Link-state routing
 - Determining link-states
 - Propagating link-states
 - Computing shortest paths
 - Dijkstra's Algorithm
- Distance-vector routing
- Inter-network routing
- Alternatives & Summary

Link-state Routing

- Two common implementations
 - Open Shortest Path First, OSPF, RFC2328
 - Intermediate-System Intermediate-System, IS-IS, RFC1142, RFC1195
- Three phases:
 - Determine link states (HELLO)
 - Broadcast link states (UPDATE)
 - Compute and install shortest paths

Determining Link States

- Use a *three-way handshake* across each link
 - “Is anyone there?”
 - “I can see you; can you see me?”
 - “I can see you.”
- Runs periodically to ensure link remains alive
 - Sometimes shortcut to alert “link down”
- Result?
 - Each router knows to whom it’s connected

Broadcast Link States

- Each router summarises and forwards
 - Each prefix represented as a link
- Simple? In principle, but in practice...
 - Versioning in case of delay, reordering
 - Summarization to make reliable
 - Number space wrapping for long uptimes
 - Flapping, convergence, loop detection, &c
- Result?
 - Each router eventually knows to whom others are connected, approximately

Compute & Install Shortest Paths

- Each router now has a representation of the network's current state
 - $\langle \text{originating-at}, \text{connected-to}, \text{metric} \rangle$
- Can run a standard shortest-path computation
 - Typically some form of Dijkstra's algorithm
 - Possibly optimize to minimize recomputation
 - Generates best *next hop* for each prefix
 - Mapped to specific interface

Dijkstra's Algorithm_(CLR, p.527)

```
# given graph  $G = (V, E)$ , and
# positive weight function  $w$ ,
# initialise costs and paths
initialise-single-source  $G\ s =$ 
1. foreach vertex  $v$  in  $V[G]$  do
2.    $d[v] = \text{infty}$ 
3.    $p[v] = \text{NIL}$ 
4.  $d[s] = 0$ 

# given subpaths  $s-u$  and  $s-v$ ,
# try  $s-u-v$  as alternative to  $s-v$ 
relax  $u\ v\ w =$ 
1. if  $d[v] > d[u] + w(u, v)$  then
2.    $d[v] = d[u] + w(u, v)$ 
3.    $p[v] = u$ 
```

```
# consider each node in turn from
# a priority queue, until all
# nodes tried
dijkstra  $G\ w\ s =$ 
1. initialize-single-source  $G\ s$ 
2.  $S = \{\}$ 
3.  $Q = V[G]$ 
4. while  $Q \neq \{\}$  do
5.    $u = \text{extract-min } Q$ 
6.    $S += \{u\}$ 
7. foreach vertex  $v$  in  $\text{Adj}[u]$  do
8.   relax  $u\ v\ w$ 
```

Contents

- Overview
- Link-state routing
- **Distance-vector routing**
 - Bellman-Ford Algorithm
 - Comparison
- Inter-network routing
- Alternatives & Summary

Distance-vector Routing

- Alternative is distance-vector
 - Routers co-operate in the computation itself
 - Should (eventually) converge
 - ...*if network is stable!*
- Protocol
 - Broadcast lowest cost to all known destinations
 - Forward using port via which best advert heard
- Common implementations:
 - Routing Information Protocol v2, RFC1723

Bellman-Ford (CLR, p.532)

```
# given graph  $G = (V, E)$ , and  
# positive weight function  $w$ ,  
# initialise costs and paths  
initialise-single-source  $G\ s =$   
1. foreach vertex  $v$  in  $V[G]$  do  
2.      $d[v] = \text{infty}$   
3.      $p[v] = \text{NIL}$   
4.  $d[s] = 0$ 
```

```
# given subpaths  $s-u$  and  $s-v$ ,  
# try  $s-u-v$  as alternative to  $s-v$   
relax  $u\ v\ w =$   
1. if  $d[v] > d[u] + w(u, v)$  then  
2.      $d[v] = d[u] + w(u, v)$   
3.      $p[v] = u$ 
```

```
# repeatedly pass over the graph,  
# relaxing each edge per node.  
# distributed version has  
# nodes doing this in parallel.  
bellman-ford  $G\ w\ s =$   
1. initialize-single-source  $G\ s$   
2. for  $i = 1 \dots |V[G]| - 1$  do  
3.     foreach edge  $(u, v)$  in  $E[G]$  do  
4.         relax  $u\ v\ w$ 
```

Comparison

- Centralized vs. Distributed computation
- State scales with #links vs. #nodes (dests)
- Network dynamics
 - E.g., Link/router failure
 - Timer and timeout management
 - DV: *count-to-infinity*
 - LS: incremental recomputation
- Management, configuration overheads
 - Easier to see what's happening with LS
 - Need to explicitly configure link weights
 - Example default: proportional to bandwidth and latency

Contents

- Overview
- Link-state routing
- Distance-vector routing
- **Inter-network routing**
 - BGP v4 – it's the only choice!
 - Route distribution and selection
 - Operations
 - Network interconnection
- Alternatives & Summary

Inter-network Routing

- An important distinction: local vs. global
 - Interior vs. Exterior Gateway Protocol (IGP, EGP)
 - Why is this important? Two reasons:
- Dynamics
 - Need to scope information propagation
- Protection (Information hiding)
 - Competition: your goals are not your neighbours'

There Can Be Only One

- Border Gateway Protocol, v4 (BGPv4)
 - Essentially distance-vector with knobs on
 - Another layer: the Autonomous System (AS)
 - Purely administrative: not relevant to data-plane
- Distance is defined as the ASPATH
 - So-called *path vector*
 - But there are many other attributes to consider
- Purpose is to enable *policy* to be applied
 - No universal (trusted) metric available

BGPv4

- Protocol for exchanging prefixes with attributes
 - Uses TCP as transport (for recursion, see recursion?)
 - OPEN, UPDATE, KEEPALIVE, (NOTIFICATION)
- OPEN sets up *sessions* between *peers*
 - Perform simple capability negotiation
 - iBGP vs. eBGP: do src and dst ASNs differ?
- UPDATES indicate
 - Withdrawn routes
 - (Shared) attributes
 - Advertised routes (Network Layer Reachability Information)

Tables, Tables, Tables

- BGP speaker typically has many sessions
 - 10? 20? 400?
- Logically maintains *Adj-RIB-In*, *-Out* for each
 - Advertisements received and to be sent
- Selection process generates *Loc-RIB*
 - Based on reachability, attributes (local-pref, aspath)
 - Resolved into per-port forwarding tables

Operations

- Scalability is a *vital* consideration
 - 300,000 prefixes, x2 per session
 - Bind to lo0 to avoid dropping all tables on link failure
 - *Default-free*: every router can handle every prefix
- Distribute internally via iBGP rather than IGP
 - Can control the dynamics much better
 - But a large network has 100s of routers!
- Route reflectors, AS confederations
 - Tweak route selection rules somewhat
- *Anycast* (1:1-of-N)
 - Advertise same *prefix* in many places. Carefully.

Network Interconnection

- How does this all fit together?
 - Roughly hierarchical (this is changing)
 - Tier-1/core/backbone vs. the rest
- Multi-homing is often desirable
 - Note that this is all *logical* though: physical diversity
- Networks interconnect via eBGP sessions
 - Points-of-Presence (Sprint, AT&T, ... customers)
 - Internet eXchanges (mutual peering)
- As ever, business and politics
 - E.g., Level3 vs. Cogent depeering

Contents

- Overview
- Link-state routing
- Distance-vector routing
- Inter-network routing
- **Alternatives & Summary**

Alternatives

- Source routing
 - But how does the host know the topology?
 - How can the network trust the host?
- Location-based routing
 - Extensive use of information embedded in environment
 - E.g., lat-long and Euclidean distance
- Map-based
 - Alternative aggregation technique for LS
- Alternative, more complex, metrics
 - Cf. QoS, later

Summary

- Routing is the process of building up information to enable efficient forwarding
- Networks are *dynamic* which makes it hard
- The two main algorithmic approaches are *link-state* and *distance-vector*
- Another operational distinction is interior vs. exterior
- BGP v4 is the only inter-domain routing protocol that counts