

Security

G54ACC – IP and Up

Lecture 9

Recap

- Authentication vs. Identification vs. Encryption
- Use of something you know/have/are
- Have previously encountered middleboxes, e.g., firewalls
- Network security involves techniques such as traffic analysis, anomaly detection
- General principles of one time pads, symmetric/asymmetric encryption, hashing

Contents

- Common Requirements
- OpenID vs. OAuth
- OpenID
- OAuth
- SSL/TLS/HTTPS

Contents

- Common Requirements
 - Network security
 - Security of the network
 - Techniques
- OpenID vs. OAuth
- OpenID
- OAuth
- SSL/TLS/HTTPS

Common Requirements

- Identification: naming a principal
- Authentication: proving your identity
 - Three-factor: know, possess, are (biometrics)
 - E.g., chip and PIN and retina/fingerprint scan
 - Two-factor: know, possess (far more common)
 - E.g., chip and PIN
- Authorization: proving you're allowed
- Confidentiality: hiding what you're doing
- Defence in depth: mitigation at all levels
 - Perfect security generally too expensive

Network Security

- Identification
 - Remote party states who they are
- Authentication
 - Prove remote party's identity
- Authorization
 - Control (authenticated) access to service
- Integrity
 - Prove a message has not been tampered with
- Confidentiality
 - Prevent third-party observation of sensitive data

Security of the Network

- An inherent conflict!
 - Controlling use of your network (but by what?)
 - How to authorise the anonymous, unknown user?
- Intrusion detection
 - Detect if host (or router) is hacked
- (Distributed) Denial of Service, BotNets
 - Handle many hacked hosts used to overload service
 - May need to distinguish from transient popularity...

Techniques

- Encryption
 - One time pads
 - Private keys (symmetric – same key to encrypt/decrypt)
 - Public keys (asymmetric – different encrypt/decrypt keys)
- Hashing: result of a one-way function
 - $\text{MAC}_k(M) = h(k, M)$
 - $\text{HMAC}_k(M) = h(k \oplus A, h(k \oplus B, M))$
 - Setting $A, B = 0x36, 0x5C$ makes collision finding harder
- Traffic analysis
 - Communication patterns may give you away
 - E.g., Hacked hosts often start using network oddly

Contents

- Common Requirements
- **OpenID vs. OAuth**
- OpenID
- OAuth
- SSL/TLS/HTTPS

OpenID vs. OAuth

- OpenID, <http://openid.net/>
 - Allows you to use one identity with multiple sites
- OAuth, <http://oauth.net/>
 - Authenticates third-party access to your data

Contents

- Common Requirements
- OpenID vs. OAuth
- OpenID
 - Dumb mode
 - Smart mode
- OAuth
- SSL/TLS/HTTPS

OpenID

- Identity
 - Not trust. Not authentication. Not authorization.
 - Enables a site to use a third-party to verify identity of a user
- Dumb (stateless) vs. Smart (stateful) modes
 - Simpler code on the consumer side
 - More computational and network resources used
 - (Optimization on dumb mode)
 - <http://wiki.openid.net/Introduction>

Dumb Mode

- Actors
 - Alice (user/you), Bob (relying party/Slashdot), Carol (provider/myOpenID.com)
- Phase 1:
 - Alice types in her identity URL with Carol to Bob's site
 - Bob GETs Carol's server URL from that page
 - Bob redirects Alice to Carol's URL adding params
 - identity, return_to, nonce
- Phase 2:
 - Carol verifies Alice *is* Alice. Somehow.
 - Carol sends Alice back adding params
 - assoc_handle (opaque handle), sig (base64 HMAC signature)
- Phase 3:
 - Bob contacts Carol POSTing all the parameters so far
 - Carol computes sig' using the secret pointed to be assoc_handle
 - If sig' == sig then tell Bob Alice *is* the Alice Carol thinks she is. Else fail.

Smart Mode

- Same as dumb mode *except*
 - Upon the first invocation, Bob and Carol setup a *shared secret*
 - In phase 2, Carol uses the shared secret as before
 - But now, in phase 3, Bob can do the check himself
- Reduces network latencies
- Reduces work done by Carol
- Shared secrets usually have a limited lifetime

Contents

- Common Requirements
- OpenID vs. OAuth
- OpenID
- **OAuth**
 - Basic flow
 - Evolution
- SSL/TLS/HTTPS

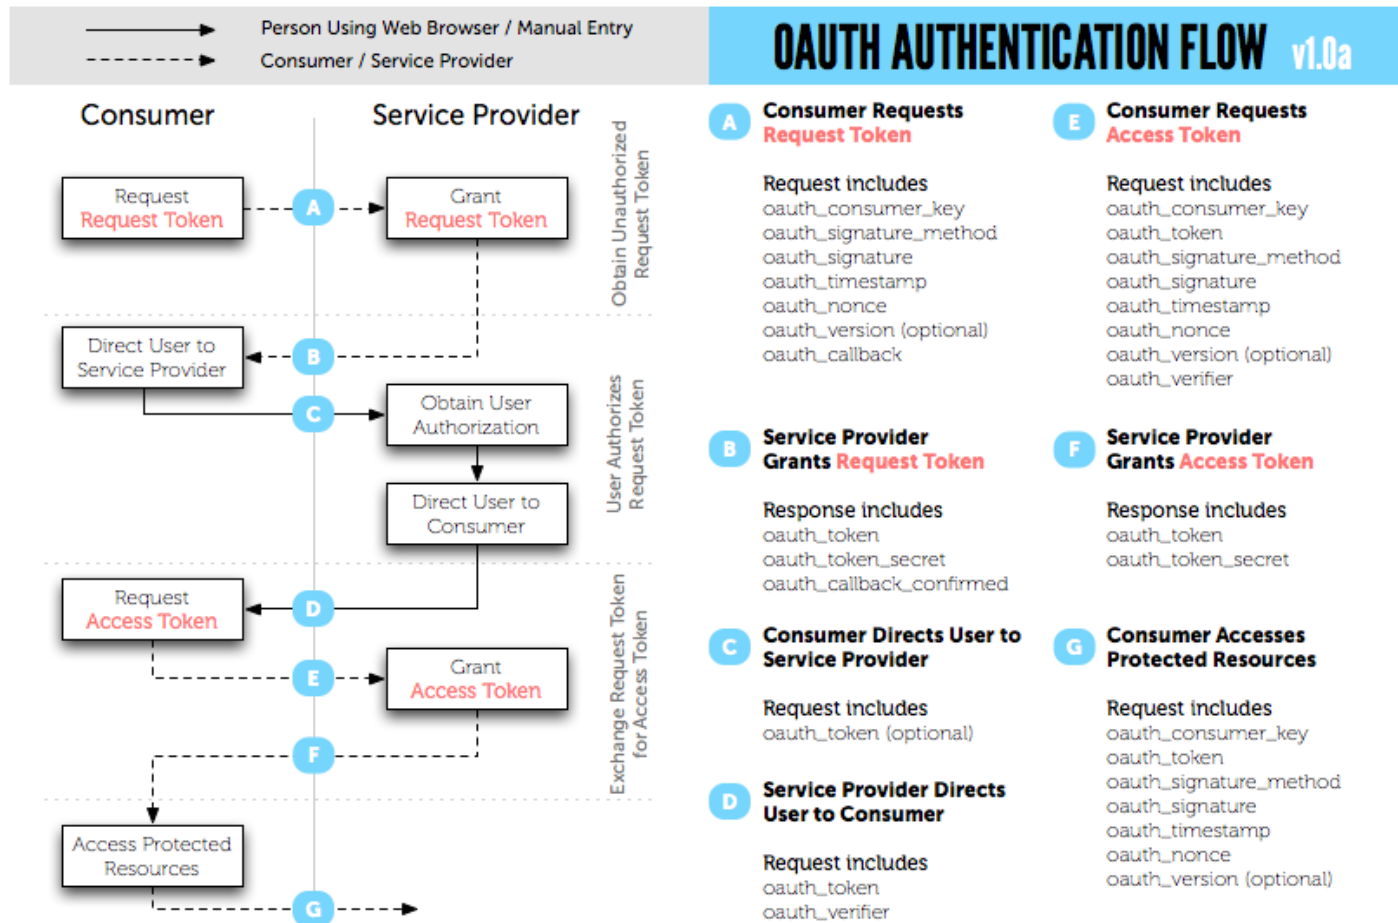
OAuth

- Appears similar to OpenID on the surface
 - In fact, it (largely) grew out of that community
 - But it serves a quite different purpose!
- Concerned with *delegating access to resources*
 - E.g., Allowing third-party apps to use your Flickr, Twitter, ... accounts
 - ...*without* giving them access to your account credentials
- <http://hueniverse.com/oauth/>

OAuth 1.0a

- Actors
 - Client (consumer/tinychat.com), server (service provider/Twitter), resource owner (user/you)
- Credentials
 - Temporary credentials (request token)
 - Token credentials (access token)
- Basic technique: HMAC-SHA1
 - Hashing incorporating a shared secret (password)
 - Prevents need to throw password around
- Additional protection via
 - Nonce (number-used-once) – but get expensive to track
 - Timestamp – enable old requests to be dropped

<http://oauth.net/core/1.0/>



OAuth 2.0

- Undergoing ratification still
- Attempts to fix problems
 - Performance at scale
 - Absence of cryptography-free options
 - Lifetime of tokens *vs.* authorization
 - Limited number of *flows*
- Adds influence from Facebook Connect
 - ...to original Flickr API Auth and Google AuthSub

Contents

- Common Requirements
- OpenID vs. OAuth
- OpenID
- OAuth
- **SSL/TLS/HTTPS**
 - Basic flow

SSL/TLS/HTTPS

- Transport Layer Security
 - Grew out of Secure Sockets Layer by Netscape
 - Incredibly complex
 - Incredibly hard to get right
 - Incredibly widely used...
- HTTPS uses TLS/SSL to
 - Provide secure channel over an insecure network
 - Verify the identity of the server
 - (Occasionally) Verify identity of the client

Using HTTPS

- For the user:
 - URLs begin https:// (TCP/443) instead of http:// (TCP/80)
- On the server
 - Generate a private/public key pair
 - Have the public key signed (signature appended using *certification authority* private key)
 - Return public certificate to client on request
- On the client
 - Client generates session key
 - Encrypts using public key to send to server
- Communication continues using symmetric encryption

Summary

- Security is complex, involving *engineering*, not just cryptography
- Generally reuse other implementations
 - Make sure you use them correctly though!
- It is a cost-benefit trade-off
 - No such thing as “perfect”
 - Sometimes using the easily available beats using the best known solution
- The system evolves: it’s an arms race