

Abstractions & 3G Evolution

G54ACC – IP and Up

Lecture 10

Recap

- Addressing, encapsulation, routing
- Multiplexing, UDP, TCP, reliability
- Address & name translation, layer violation
- XMPP, HTTP, pipelining, P2P, active networks
- IP network management
- Sockets, multiplexing (again), endianness
- Security, OpenID, OAuth, network security

Contents

- System Design
- Common Techniques
- 3G Data Network Evolution

Contents

- System Design
 - Goals
 - Resources
 - Constraints
- Common Techniques
- 3G Data Network Evolution

System Design

- The art and science of extracting what you can from what you have
- Specifically,
 - Solve the problem,
 - Using your resources,
 - Subject to any constraints

Goals and Resources

- Goals are often not well-formed
 - Scalability, modularity, extensibility, elegance
- Resources are a multiplicity of
 - Computation
 - Storage
 - Communications
 - (People?)

System Constraints

- Design
 - Bottleneck resource
 - The most constrained element
 - Scaling
 - To infinity and beyond!
- Resources
 - Latency
 - c , the speed of light
 - Throughput, capacity
 - Money

Social Constraints

- Standards compliance
 - Legal implications
- Market conditions
 - (Economic) sustainability
- Labour availability
 - Otherwise nothing gets built
- Future proofing
 - Times change

Contents

- System Design
- Common Techniques
 - Layering & Tunneling; Extensibility
 - Hierarchy; Binding & Indirection
 - State management; Randomization
 - Damping & Hysteresis; Operating timescale
- 3G Data Network Evolution

Layering & Tunneling

- Control complexity by restricting component interactions
 - Simplicity at the cost of performance
 - Often benefit from (restricted) layer violation
- Need not always be strictly layered
 - Protocol heaps [HOTNETS'02]
 - Cycles due to tunneling
- E.g., IP and OSI stacks; IP-in-IP

Extensibility

- “Prediction is very difficult, especially about the future” (Niels Bohr)
- You don’t know how things will be used
 - Or in what circumstance
- Allow for future growth & upgrade
 - “Ships in the night” vs. integrated operation
 - Tunneling as a workaround
- E.g., Version fields, negotiation, explicit state

Hierarchy

- Recursive decomposition into smaller pieces
 - Each typically depends on parent
- No single point of control
 - Though potential single point of failure
- Inter-sibling communication can be expensive
 - May choose to introduce shortcuts (complexity)
- Structure may become too rigid
 - Need to adapt to system dynamics
 - Cf. Pooling, P2P
- E.g., CIDR, IANA

Binding & Indirection

- Need to translate from abstraction to instance
 - Computation vs. lookup
 - Distributed vs. centralised control
- If well-known, translate (bind) automatically
 - Choice of when: late vs. early binding
- E.g., DNS, service lookup
- Combine with multiplexing: *virtualization*

State Management

- Soft state
 - Requires refreshing else deleted on timer
 - Increased bandwidth for automatic clean up after failure
- Explicit state
 - Easier to interpret, manage, error-protect
- E.g., Link state vs. distance vector routing

Randomization

- A powerful tool
 - Tie-breaking
 - Security
- Used up and down the stack
- E.g., Ethernet collisions, nonces, designated router election

Damping & Hysteresis

- (Reducing amplitude vs. adding memory)
- System state changes based on thresholds so may fluctuate if near a threshold
- Such fluctuations, if rapid, are usually bad
 - Bandwidth overheads
 - Reduced performance
 - Increased likelihood of inconsistency
- Make thresholds *state-dependent*
- E.g., TCP congestion control/RTT estimation; but BGP flap damping is problematic

Operating Timescales

- Not all operations need be equally fast
- Not all operations have same granularity
 - E.g., Once per packet vs. once per communication
- Data vs. control vs. management
- E.g., Connection setup, QoS for packets vs. flows vs. aggregates

Pipelining

- A complex task can be subdivided
 - Into many independent tasks: fully parallelizable
 - Into many dependent tasks: not much help
- Many *serially dependent* subtasks
 - Each depends only on predecessor
 - Typically improves throughput
 - ... while slightly increasing latency
- E.g., compare HTTP/1.1 vs. HTTP/1.0

Batching & Locality

- Grouping tasks to *amortize* overheads
 - Assumes $(\text{overhead for } N) < N * (\text{overhead for } 1)$
 - Increased latency accumulating batch
vs. higher throughput
- Recently used data probably used again, soon
 - Basis for use of caching
 - Increase overheads for statistical benefits

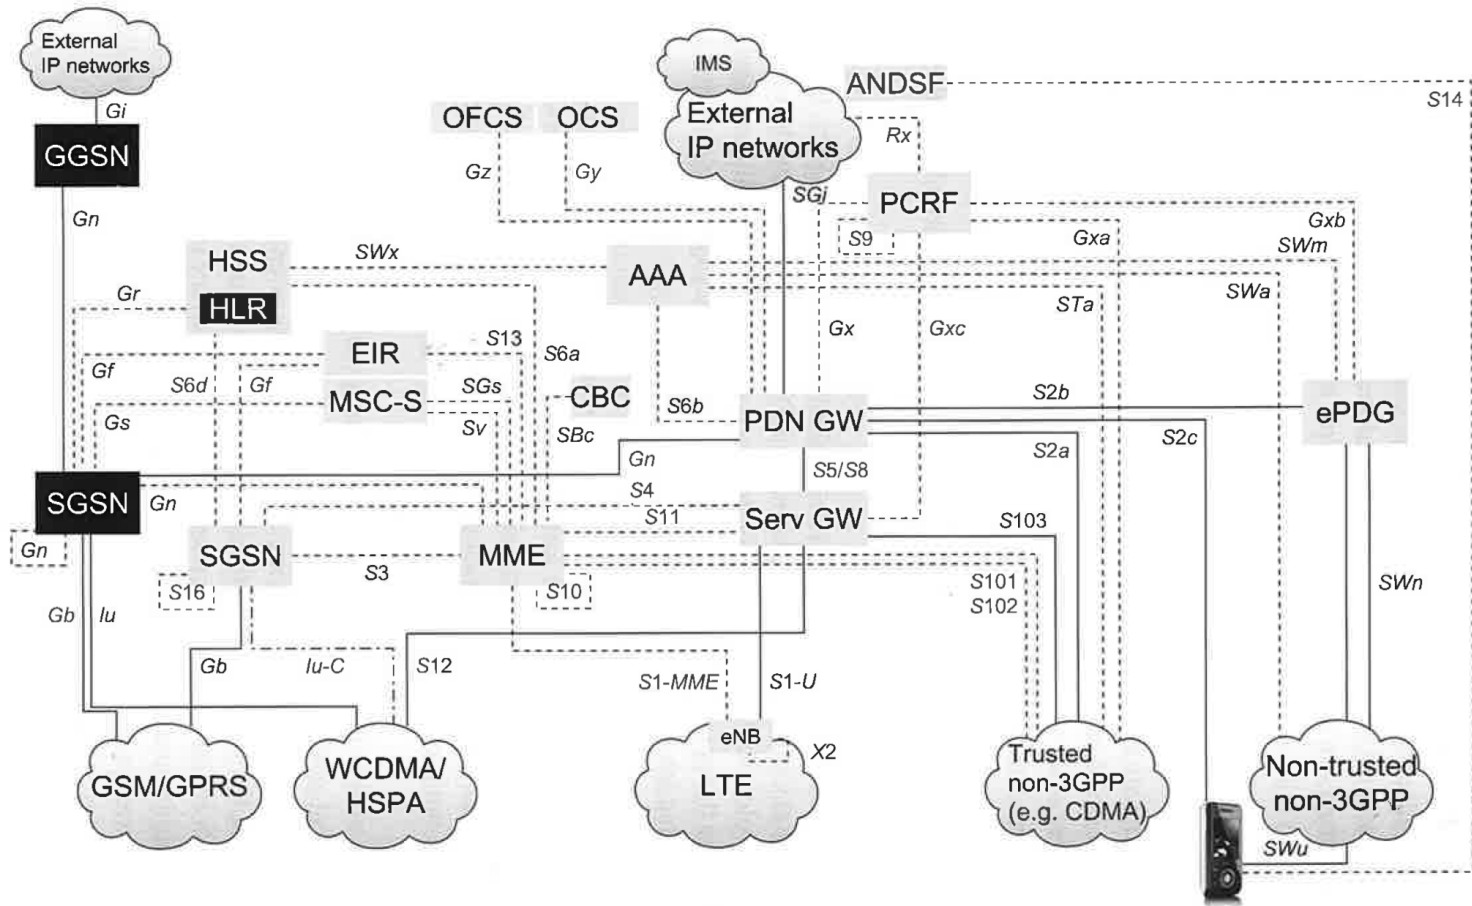
Contents

- System Design
- Common Techniques
- 3G Data Network Evolution
 - Basic LTE operation
 - A different style
 - Some interesting points

EPC Data Networks

- Vision is to provide continuous IP connectivity
- ...on resource constrained devices
 - Single radio, battery powered
- ...while mobile
 - Geographically, at speed
 - Logically, between network types (2g, 3g, WLAN)
- ...with detailed control/management facility
 - Policy
 - Charging
 - Lawful intercept

EPC Data Networks



From Olsson *et al.*, Elsevier AP 2009,
 “SAE and the Evolved Packet Core:
 Driving the Mobile Broadband Revolution”

Basic LTE Operation

- Tx
 - UE registers with MME
 - Obtains configuration (address, Serv GW, &c)
 - Network signals bearers (tunnels) to be set up
- Rx
 - MME pages Serv GWs for UE
 - UE responds, setting up bearers via Serv GW – PDN GW

A Different Style

- It's the network, stupid!
 - Not the stupid network
 - Cf. end-to-end argument and the “dumb network”
 - *Dramatically* more complex accounting required
- E.g.,
 - PCRF, OFCS, OCS
 - Apply complex connectivity and charging policies
 - EIR
 - Deny connectivity to stolen devices
 - ETWS
 - Earthquake and Tsunami Warning Service
 - (to be replaced by Public Warning Service, to benefit US/EU)

Signal Measurement

- eNodeBs (base stations, eNBs) may overlap
 - Which should I pick? How?
- eNB creates idle slots to enable UEs to measure signal strength
 - Only rely on a single radio
 - Can't do simultaneous Rx and Tx
 - Can't be simultaneously on different RANs
- UE switches as desired

Idle Mode

- UE can go into idle mode to save power
 - Less background transmission to network
 - But what about tracking location, e.g., to deliver incoming packets?
- Requires support for paging by MME
 - At request of Serv GW, while it buffers packets
- Periodic updates support vanished terminals
 - Else might keep paging forever
- Trade-off paging overheads vs. UE battery

Tracking Areas

- To where should the page be sent?
 - Each UE gets a list of eNBs defining TA
 - UE notifies network if it changes TA
- What about, e.g., trains though?
 - Massive synchronized movement between TAs
 - E.g., Bearer setup/teardown
- Solve by giving out different TA lists
 - Spreads the load
- Further optimize using movement prediction

Pooling

- How to dimension capacity?
 - 2G/3G used hierarchical core
 - UE in a cell connects to 1 base station, &c.
- Region in EPC has a *pool* of weighted MMEs
 - Enables eNBs to distribute load across MMEs
- More efficient use of statistical multiplexing; compare:
 - Peak + 20% everywhere, e.g., region of 1M users
 - Pool of $N+1$ nodes where each node handles X users

Summary

- System design is hard: an art and a science
- There are many *many* details and interactions involved in any large-scale (networked) system
 - All networked systems are thus interesting at some level
- The Internet and related technologies are good widely deployed, evolved, examples
 - Demonstrate a large toolkit of useful techniques and abstractions
 - Occasionally (partly) by accident rather than design