

Reliability

G54ACC

Lecture 11

richard.mortier@nottingham.ac.uk

Recap

- Congestion control cf. flow control
- Adapt transmission window based on timers
- TCP reacts to loss by backing off
 - Loss is assumed to be caused by network congestion
- TCP reacts to successful delivery by probing for more bandwidth
 - In the simple case, generates a “saw tooth” usage

Contents

- Performance and reliability
- TCP congestion control
- Multimedia

Contents


- Performance and reliability
 - Simple schemes
 - Congestion collapse
- TCP congestion control
- Multimedia

Achieving Reliability

- Reconstruct lost data
 - E.g., Forward Error Correction
 - Cf. coding and block codes
- Retransmit lost data
 - ...but how to detect loss?
- Rx explicitly acknowledges: ARQ
 - Automatic Repeat ReQuest
 - Alternative: negative acknowledgements
 - Cf. OSPF/ISIS

Stop 'n' Wait

- Simplest possible
 - Tx seq(x)
 - Tx wait for ack(x)
 - Tx seq(x+1)
- Really poor performance in high-latency high-bandwidth networks
 - Why?




Reliability and Performance

- *Rate control*: never send too fast for network
 - How to estimate appropriate rate?
- *Sliding window*: allow unacknowledged data in flight
 - How to determine correct window size?
- *Retransmission Timeout*
 - How long to wait to decide a segment is lost?
- Require estimates of dynamic quantities
 - Actually, proxies for the network's current capacity

First Attempt (pre-1988)

- Go-Back-N ARQ
 - Permit N segments in flight
 - Timeout implies loss
 - Retransmit from lost packet onward
- Self-clocking behaviour
 - Tx next segment after previous segment acknowledged
- Loss wastes throughput
 - *Goodput vs. throughput*



Congestion Collapse

- If network load gets too high, experience *congestion collapse*
 - Initially observed on NSFNet Oct/1986
 - LBNL—UCB (400yds/3 hops) dropped 32kb/s to 40b/s
- Why?
 - Router buffers fill, traffic is discarded, hosts retransmit
 - But early TCP (4.2BSD) was naive! Retransmit **more** since data was lost...
 - Boom!
- Solution: Van Jacobson's congestion control algorithm (SIGCOMM'88): Slow-start and Congestion avoidance

Contents

- Performance and reliability
- TCP congestion control
 - Avoidance and recovery
 - Tahoe, Reno, NewReno, SACK, Vegas
 - ECN marking
- Multimedia

Congestion Control

- Need to determine current parameters
 - Window size for most TCPs
 - Either estimate, or *actively probe*
- Additive-Increase, Multiplicative-Decrease
 - Maintain a *congestion window*, cwnd
 - awnd gives Rx window flow control
 - $wnd = \min(awnd, cwnd)$
- First introduced as TCP Tahoe in 4.3BSD (1988)

TCP Tahoe

- Congestion Avoidance: given $wnd = N$ segments
 - Segment acknowledged: $cwnd += 1/cwnd$
 - RTO: $cwnd := 1$; $ssthresh /= 2$
 - But how to initialise $cwnd$?
- Slow-Start (actually, *exponential growth!*)
 - Original behaviour was constant initial $cwnd$
 - Every segment acknowledged, $cwnd += 1$
 - Stops when loss occurs or $cwnd == ssthresh$
 - Adjust $ssthresh$ up once we reach avoidance

Fast Retransmit

- Detecting just a single loss takes a full RTO
 - 500ms or more, at 500ms granularity
- Rx can signal a “hole” in the stream
- Send *duplicate ACKs (dup-acks)*, typically 3
- Tx retransmits the relevant segment
- $ssthresh := \max(wnd/2, 2);$
 $cwnd := 1$
- Return to slow-start

Evolution: Reno, NewReno, SACK

- Reno (4.4BSD, 1990) adds *fast recovery*
 - Dup ACKs trigger retransmit
 - But now set $cwnd \neq 2$
 - Then inflate wnd by number of dup ACKs as self-clocking means segments have left the pipe
- NewReno (4.4BSD, 1993) modifies fast recovery
 - Reno has poor performance for multiple losses
 - ACK for part of window no longer exits recovery
- Selective ACKnowledgement
 - Maintain more state to retransmit *only* lost segments

Revolution: Vegas

- Alternative approach to congestion avoidance
 - Source uses RTT to estimate #packets in pipe
 - End-to-end queuing delay as congestion measure
- for every RTT:
 - $RTT_{min} = \min(RTT, RTT_{min})$
 - if $wnd/RTT_{min} - wnd/RTT < \alpha$ then $wnd++$
 - if $wnd/RTT_{min} - wnd/RTT > \beta$ then $wnd--$
- Never widely deployed

Marking

- Many *many* other variants
 - FACK, VENO, BIC, CUBIC, HYBLA, Compound, Westwood, Fusion, ...
 - Cf. doi:10.1109/SURV.2010.042710.00114
- Basically: loss signals congestion, then react
 - But waiting for loss takes a full RTO
- Explicit Congestion Notification, RFC3168
 - Signal congestion *early* via flags
 - Requires router support: marking, queuing

Contents

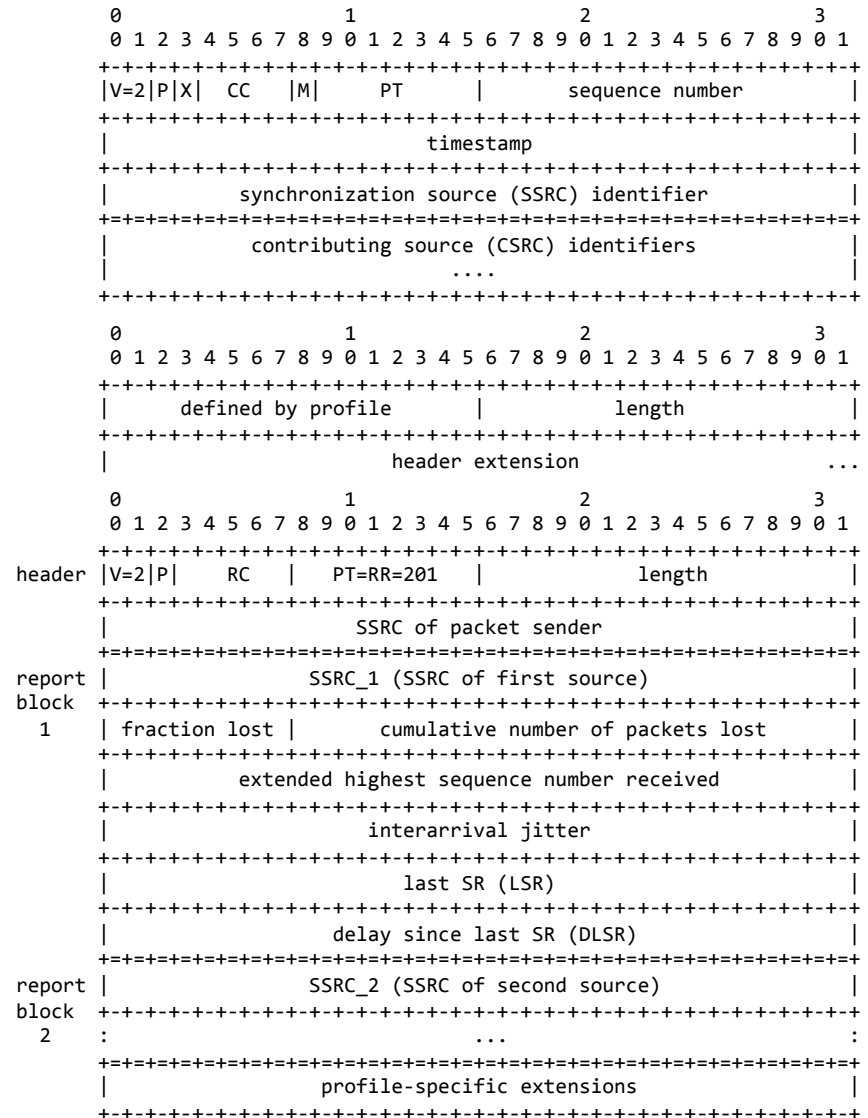
- Performance and reliability
- TCP congestion control
- **Multimedia**
 - $RT \cdot P$
 - TCP friendliness
 - SST

Multimedia

- TCP is not always useful
 - Messages/frames preferably to bytestream
 - Reliability can cause untimely delivery
 - Canonical example: Real-time media
- Audio/video codecs usually produce frames
 - Not a continuous bytestream
- Losing a frame is better than delaying all subsequent data
 - We can skip or interpolate missing frames

Real Time Protocol

- Encapsulates media over UDP
 - Sequencing, timestamping, delivery monitoring
 - No reliability, no QoS
- Adds a control channel (RTCP)
 - Backchannel to report statistics, participants, &c
- Transport only
 - Leaves encodings, floor control, &c to application
- Extensible



TCP Friendliness

- Real-time media being sent at a given rate
 - May cause TCP to experience loss and backoff
 - Media will eventually win
- Need to adapt real-time media
 - Equation based congestion control
 - $T = \text{avg}(\text{wnd}) / \text{RTT}$
 - $p(\text{drop}) = 2 / (3 * w^2) \Rightarrow w = \text{sqrt}(2 / 3p)$
 - Thus $T(p) \approx (1 / \text{RTT}) * \text{sqrt}(3 / 2p)$
- Send drop rate to Tx and it uses $T(p)$ directly

Structured Stream Transport

- Presented in SIGCOMM'07
- Single TCP stream imposes total order
 - Cf. HTTP/1.1 later
- Alternative: enable “substreams” to be created
 - More flexible scheduling by both ends
 - Treat datagram transmission as ephemeral substream
 - Multiplex substreams onto congestion controlled channels

Contents

- Performance and reliability
 - Simple schemes
 - Congestion collapse
- TCP congestion control
 - Avoidance and recovery
 - Tahoe, Reno, NewReno, SACK, Vegas
 - ECN marking
- Multimedia
 - RT*P
 - TCP friendliness
 - SST

Summary

- Achieving both reliability and performance is *hard!*
- There're many subtle interactions particularly if your scheme is to be general
 - Even trickier if you mix in non-responsive traffic
- But try not to let TCP become a strait-jacket...

Quiz

1. What are the different tradeoffs made by FEC and ARQ in providing reliability?
2. How does each behave on a very reliable network?
3. How does each behave on a very unreliable network?
4. Why does stop'n'wait perform so badly on a high-bandwidth, high-latency network?
5. Sketch the evolution of the window size for TCP Tahoe and TCP Reno, assuming no loss until **after** exiting slow-start.
6. Why does TCP Tahoe suffer in the face of a burst of losses? Why was this particularly bad?
7. Why might you **not** want TCP's features when transporting real-time multimedia data?