

Programming

G54ACC – IP and Up

Lecture 8

Recap

- Basics of Berkeley Sockets API
 - `socket()`, `connect()`, `send()`, `recv()`, &c
- Endianness
- Multiplexing at lower layers
 - We'll discuss multiplexing from the programmer point of view
- Remote Procedure Call (RPC)
 - Invocation of code on a remote machine

Contents

- Berkeley Sockets
- Web RPC
- RESTful web services
- Flavours of “cloud computing”
- Summary

Contents

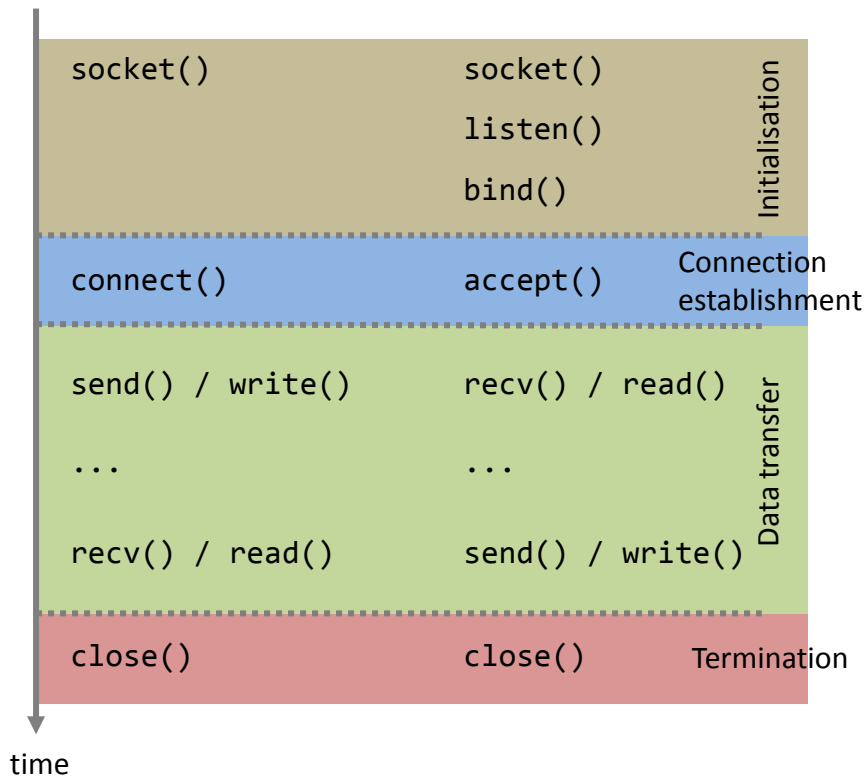
- Berkeley Sockets
 - TCP/UDP
 - Endianness
 - Multiplexing
- Web RPC
- RESTful web services
- Flavours of “cloud computing”
- Summary

Sockets

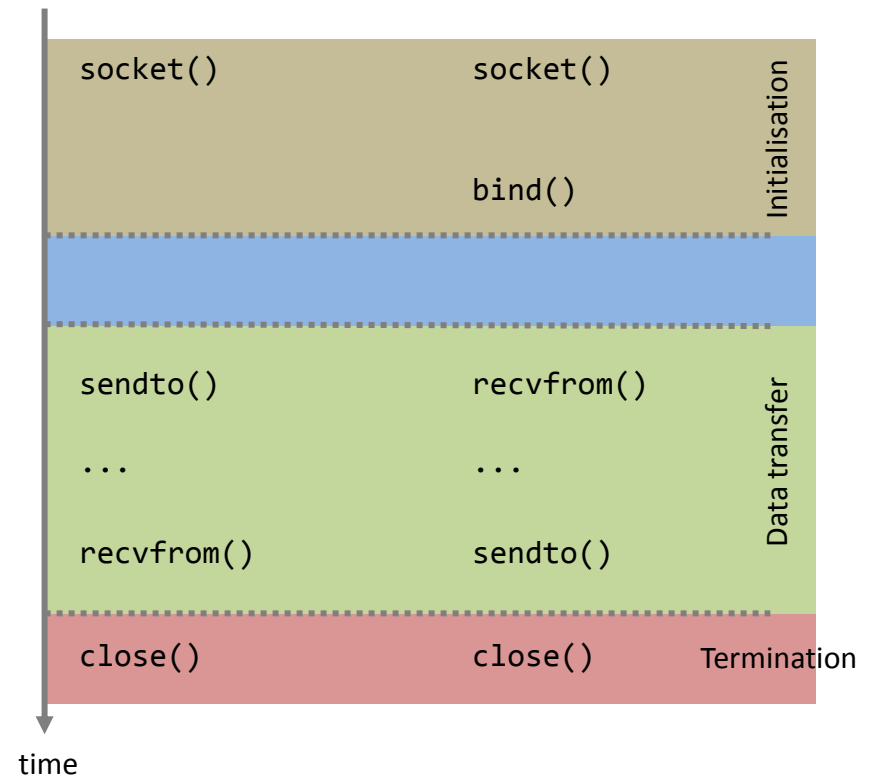
- Berkeley Sockets API, 4.2BSD (1983)
 - Use *Internet sockets*, presented as file descriptors
 - Extended to many other protocols
- `int socket(int domain, int type, int protocol)`
 - domain: communication domain
 - `PF_INET`, `PF_INET6`, `PF_LOCAL`, `PF_RAW`
 - type: communication semantics
 - `SOCK_STREAM`, `SOCK_DGRAM`, `SOCK_RAW`
 - protocol: usu. 1:1 with (domain,type)
 - `IPPROTO_ICMP`, `IPPROTO_TCP`, `IPPROTO_UDP`

TCP vs. UDP

- File-like interface
 - Streaming



- Packet-like interface
 - Need to packetise



Endianness

- Host byte order need not be identical
 - Big endian (most-significant-byte first)
 - vs. Little endian (least-significant-byte first)
- Need to decide which we use on the network
 - Lest 0x0001 be received as 0x0100
 - Or negotiate every time?!
- Network byte order == big endian
 - No reason. It just is.
 - Convert via `u[16|32] htonl[s|1]/ntoh[s|1](u[16|32])`

Multiplexing

- Naive: read and block
 - But what if you need to handle >1 socket?
- Naive 2: poll non-blocking socket
 - Wastes CPU
- Use `select()`, either blocking or non-blocking
 - Enables waiting (until timeout) on any socket of a set becoming active
 - ```
int select(int nfd, fd_set *rfd, fd_set *wfd, fd_set *efd,
 struct timeval timeout)
```
- Check W. Richard Stevens “*Unix Network Programming*” and “*TCP/IP vol.1*” for details



# Contents

- Berkeley Sockets
- Web RPC
  - XML-RPC
  - SOAP
- RESTful web services
- Flavours of “cloud computing”
- Summary

# Web RPC Systems

- (At least) Two: XML-RPC *vs.* SOAP
  - Appears to be/have been a religious war here
  - SOAP evolved from XML-RPC
- Both provide
  - Data encoded as XML
  - Remote Procedure Call interfaces
  - ...to remote web services

# XML-RPC

- Older, simpler
  - <http://www.xmlrpc.com/spec>
- Community based

```
POST /RPC2 HTTP/1.0
User-Agent: Frontier/5.1.2 (WinNT)
Host: betty.userland.com
Content-Type: text/xml
Content-length: 181
```

```
<?xml version="1.0"?>
<methodCall>
 <methodName>examples.getBounds</methodName>
 <params>
 <param>
 <value><i4>41</i4></value>
 </param>
 </params>
</methodCall>
```

```
HTTP/1.1 200 OK
Connection: close
Content-Length: 158
Content-Type: text/xml
Date: Fri, 17 Jul 1998 19:55:08 GMT
Server: UserLand Frontier/5.1.2-WinNT
```

```
<?xml version="1.0"?>
<methodResponse>
 <params>
 <param>
 <struct>
 <member>
 <name>lowerBound</name>
 <value><i4>18</i4></value>
 </member>
 <member>
 <name>upperBound</name>
 <value><i4>139</i4></value>
 </member>
 </struct>
 </param>
 </params>
</methodResponse>
```

# SOAP

- Newer, more featureful, more flexible
  - Doesn't require HTTP as transport
  - Provides e.g., discovery (WSDL), &c
- Not particularly simple
- W3C standard

```
POST /InStock HTTP/1.1
Host: www.example.org
Content-Type: application/soap+xml;
charset=utf-8
Content-Length: nnn

<?xml version="1.0"?>
<soap:Envelope
xmlns:soap="http://www.w3.org/2001/12/s
oap-envelope"
soap:encodingStyle="http://www.w3.org/2
001/12/soap-encoding">

 <soap:Body
xmlns:m="http://www.example.org/stock">
 <m:GetStockPrice>
 <m:StockName>IBM</m:StockName>
 </m:GetStockPrice>
 </soap:Body>

</soap:Envelope>
```

# Contents

- Berkeley Sockets
- Web RPC
- RESTful web services
  - REST constraints and interfaces
  - Creating a RESTful service
- Flavours of “cloud computing”
- Summary

# RESTful Web Service

- Roy Fielding during HTTP/1.1 development (1994–2002)
  - doi:10.1145/514183.514185
  - [http://www.ics.uci.edu/~fielding/pubs/dissertation/rest\\_arch\\_style.htm](http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm)
- Goals
  - Minimize latency, communication
  - Maximise implementation scalability, component independence

# RESTful Web Service

- A RESTful service conforms to the REST constraints
  - Representational State Transfer
  - Makes more extensive use of HTTP features
    - In principle can use other transports
- Client requests generate server responses
- Both built around *representations of resources*
  - HTML, JPEG, whatever in principle. JSON or XML in practice.

# REST Constraints

- Six architectural constraints
  - Client-server, Cacheable, Layered
    - Separate concerns, beware existence of caches, proxies, &c
  - Code on demand (optional)
    - Server can push e.g., JavaScript to client as part of response
  - Stateless, Uniform interface
    - ...



# Stateless

- Does *not* mean the server must be stateless!
- Simply that no *client* context is stored on the server between requests
  - Each client request contains enough to service it
  - If server does store state, it is addressable as a resource via URL in the usual way
- Contrast

GET /resources/nextPage?

...

```
<?xml version="1.0"?>
<rsp status="ok">
 <resource id="1" />
 <resource id="2" />
</rsp>
```

GET /resources/?page=2

...

```
<?xml version="1.0"?>
<rsp page="2" nextPage="3">
 <resource id="21" />
 <resource id="22" />
</rsp>
```

# Uniform Interface

- Client uses representation to modify server state, addressed via URL
- RESTful service maps CRUD to HTTP methods
  - CREATE: POST
  - READ/RETRIEVE: GET, idempotent, side-effect free
  - UPDATE: PUT, idempotent (since uses specified name)
  - DELETE: idempotent
- CREATE/UPDATE as POST vs. PUT appears debated
  - Differ in treatment of Request-URI and request entity
  - POST: treat entity as “new subordinate”
  - PUT: entity “stored under supplied Request-URI”

# Creating a RESTful Service

- What are the URIs?
  - Find the nouns, define *collections*
- What's the format?
  - XML, JSON, &c
- What methods are supported at each URI?
  - Mapping to the CRUD functions
- What status codes could be returned?
  - 1xx (informational), 2xx (successful), 3xx (redirection), 4xx (client error), 5xx (server error)

# Notable Points

- Collections
  - <http://.../people> : GET (list), POST (append), PUT (replace), DELETE
  - <http://.../people/1> : GET (retrieve), PUT (update), DELETE
- Tickets used to provide *at-most-once* semantics
  - Unique, server-generated GUIDs
  - POST to /ticket/o to create a “slot” in the server
- Redirects help clients handle tickets
  - Ticket response redirects to /o/<n>
  - PUT /o/<n> creates new item, idempotently

# Contents

- Berkeley Sockets
- Web RPC
- RESTful web services
- Flavours of “cloud computing”
- Summary

# Cloud Computing

- Datacenter as a utility
  - Buy it (rent it) when you need it (OpEx vs. CapEx)
  - Scales “on demand”
- Beware!
  - Variable performance
  - Little control over data placement
  - Unknown trust implications
  - Cloud provider as single-point-of-failure
    - Not just their hosts, but their and your network connectivity
  - Building scalable services is still hard!

# Flavours

- X as a Service for X =
  - Software: Salesforce
    - Provides you with a Customer Relationship Manager service
    - No longer need to manage software, upgrades, &c
  - Platform: Google App Engine, PiCloud, Heroku, Azure
    - Hosts web apps backed by shared store (SQL-like)
    - Fires up (Python, Java, Ruby, CLR) VM as required
  - Infrastructure: Amazon EC2, Rackspace
    - Provides XEN hosted VMs (Linux, Windows, custom)
    - Also persistent store (S3), other infrastructure APIs

# Summary

- There are a wide range of APIs that attempt hide the complexity of network programming
  - Sockets is the (old) canonical Internet API on which all else is built
  - Web RPC was briefly popular
  - RESTful platforms with APIs are “web2.0” and all the rage (and jolly useful)
- Cloud computing in some ways goes a stage further, hiding even more of the complexity
  - But you introduce other problems due to, e.g., the scale of the application
- Failures are always a pain to deal with, anywhere