

Applications

G54ACC

Lecture 15

richard.mortier@nottingham.ac.uk

Contents

- Higher layers
- HTTP
- XMPP
- And the radically different...

Contents

- Higher layers
 - Session
 - Presentation
 - Structure
- HTTP
- XMPP
- And the radically different...

Session Layer


- Open/close semi-permanent dialogue
- SSH, RFC4251 *et al*
 - Provide encryption, keepalives for connection
- RTCP, RFC3550
 - Floor control: who's talking, who's listening
- Rather forced
 - Not traditionally part of TCP/IP stack

Presentation Layer

- Conversion of data encodings
 - ASCII vs. Unicode UTF8
 - Unix (LF) vs. MS-DOS (CRLF)
 - MPEG
- Also rather forced in TCP/IP terms
 - ...and can get quite weird, cf. FAT filesystem madness on Linux
- Functions still exist, just not explicitly layered
 - Generally lumped together in “application layer”


Client/Server

- More common, easier to implement
- Can be hard to scale for many clients
- Requires always-on server
 - E.g., HTTP, XMPP



Peer to Peer

- Scales naturally but hard to manage
 - Intermittent service
 - Dynamic peers
- Can have unfortunate effects on network use
- E.g., BitTorrent, Skype




Contents

- Higher layers
- HTTP
 - Objects
 - Requests, Responses
 - State
 - Network Usage
 - Security
- XMPP
- Radically different

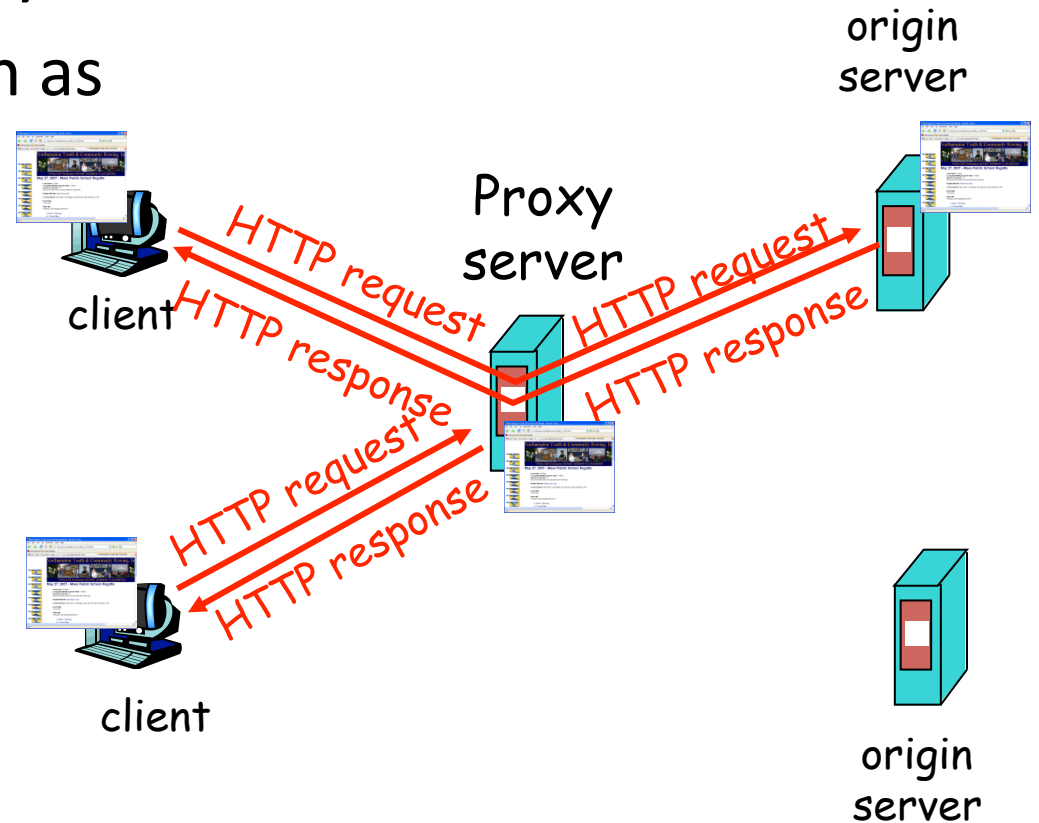
HyperText Transport Protocol, HTTP

- Client-server, request-response
 - *Objects*, named via URLs
 - *Clients* (or *user agents*), retrieve objects from ...
 - *Servers*, which store or generate objects



Can Be Made Complicated

- *Proxies*, get in the way
- Provide features such as
 - Caching,
 - Logging,
 - Transcoding, &c.



Objects

- Used to be HTML pages
 - HyperText Markup Language
- Might now be just about anything
 - Page, image, endpoint, computation
- Labelled via a Uniform Resource Identifier, URI
 - In the web, this is a Uniform Resource Locator, URL
 - `scheme://host:port/path/to/resource/?query`
 - Rarely, might be Uniform Resource Name, URN
 - `urn:ietf:rfc:3986`

Requests

- HTTP/1.0
 - Connect TCP/80
 - Issue request method
 - GET, POST, HEAD
 - Issue headers
 - Language, &c.
 - Process result
- Issue further requests as required
 - Images, &c.
 - Separate connections

```
$ telnet google.com 80
Trying 173.194.37.104...
Connected to google.com.
Escape character is '^]'.
GET / HTTP/1.0
```

```
HTTP/1.0 302 Found
Location: http://www.google.co.uk/
Cache-Control: private
Content-Type: text/html; charset=UTF-8
Set-Cookie:
  PREF=ID=76b0229e458ed281:TM=1283786147:LM=1283786147:S
  =YChkvG74Grq1F0nS; expires=Wed, 05-Sep-2012 15:15:47
  GMT; path=/; domain=.google.com
Set-Cookie: NID=38=E2-
  NE5vYotdt6QPD8ENPi0rLaI3DJUS635jvNkw8AKMIRFp37i1jV8G6j
  Pik3wvrWdMQRvw2BI1PKLp-
  WS3bhZuRZ6lHZZfgfQDqXje6gb5BIXgBxATV_N1G1h-Lkqj3;
  expires=Tue, 08-Mar-2011 15:15:47 GMT; path=/;
  domain=.google.com; HttpOnly
Date: Mon, 06 Sep 2010 15:15:47 GMT
Server: gws
Content-Length: 221
X-XSS-Protection: 1; mode=block
```

```
<HTML><HEAD><meta http-equiv="content-type"
content="text/html; charset=utf-8">
<TITLE>302 Moved</TITLE></HEAD><BODY>
<H1>302 Moved</H1>
The document has moved
<A HREF="http://www.google.co.uk/">here</A>.
</BODY></HTML>
```

Responses

```
$ telnet google.com 80
Trying 173.194.37.104...
Connected to google.com.
Escape character is '^]'.
GET / HTTP/1.0
```

```
HTTP/1.0 302 Found
Location: http://www.google.co.uk/
Cache-Control: private
Content-Type: text/html; charset=UTF-8
Set-Cookie:
PREF=ID=76b0229e458ed281:TM=1283786147:LM=1283786147:S
=YChkvG74Grq1FOnS; expires=Wed, 05-Sep-2012 15:15:47
GMT; path=/; domain=.google.com
Set-Cookie: NID=38=E2-
NE5vYotdt6QPD8ENPiOrLaI3DJUS635jvNkw8AkMIRFp37i1jV8G6j
Pik3wvrWdMQRvw2BI1PKLp-
WS3bhZuRZ6LHZZfgfQDqXje6gb5BIXgBxATV_N1GLh-Lkqj3;
expires=Tue, 08-Mar-2011 15:15:47 GMT; path=/;
domain=.google.com; HttpOnly
Date: Mon, 06 Sep 2010 15:15:47 GMT
Server: gws
Content-Length: 221
X-XSS-Protection: 1; mode=block
```

```
<HTML><HEAD><meta http-equiv="content-type"
content="text/html; charset=utf-8">
<TITLE>302 Moved</TITLE></HEAD><BODY>
<H1>302 Moved</H1>
The document has moved
<A HREF="http://www.google.co.uk/">here</A>.
</BODY></HTML>
```


- Status line
- *Response headers*
 - Provide meta-data
 - Location, type, length, &c.
- **Content**
 - Generated
 - Read from file, &c.
- **Stateless**
 - No protocol level server-side link between requests
 - Useful for unreliable links

Handling State

- Necessary for some applications
 - E.g., Shopping carts, preferences, usage tracking, &c.
- Let the client do it
 - Avoids server scaling issues
 - Needs care with authentication though
- Cookies
 - Intended to be small
 - Server headers *Set-Cookie: <cookie-value>*
 - Client can then use *Cookie: <cookie-value>* in subsequent requests


A Connection per Request

- Inefficient if objects retrieved one at a time
- TCP unfair if requests handled in parallel
 - Go through slow-start every time
 - Also somewhat inefficient due to TCP overheads
 - E.g., Old Netscape “limit to 4 connections” behaviour



Network Usage

- HTTP/1.1
 - Persistent connections
 - Thus, multiple requests on a single connection
 - Sharing/reusing TCP state
- Scheduling issues
 - I ask for A, B, C
 - You can't reply C until I receive A and B
 - Why might this be a problem?



Security

- Transport Layer Security, TLS
 - Grew out of Netscape Secure Sockets Layer, SSL
 - Incredibly complex
 - Incredibly hard to get right
 - Incredibly widely used...
- HTTPS uses TLS/SSL to
 - Provide secure channel over an insecure network
 - Verify the identity of the server
 - (Occasionally) Verify identity of the client

Using TLS in HTTPS

- For the user
 - URLs begin https:// (TCP/443) instead of http:// (TCP/80)
 - But how does the user know it means anything?
- On the server
 - Generate a private/public key pair
 - Have the public key signed (signature appended using *certification authority* private key)
 - Return public certificate to client on request
- On the client
 - Client generates session key
 - Encrypts using public key to send to server
- Communication continues, using symmetric encryption

Contents

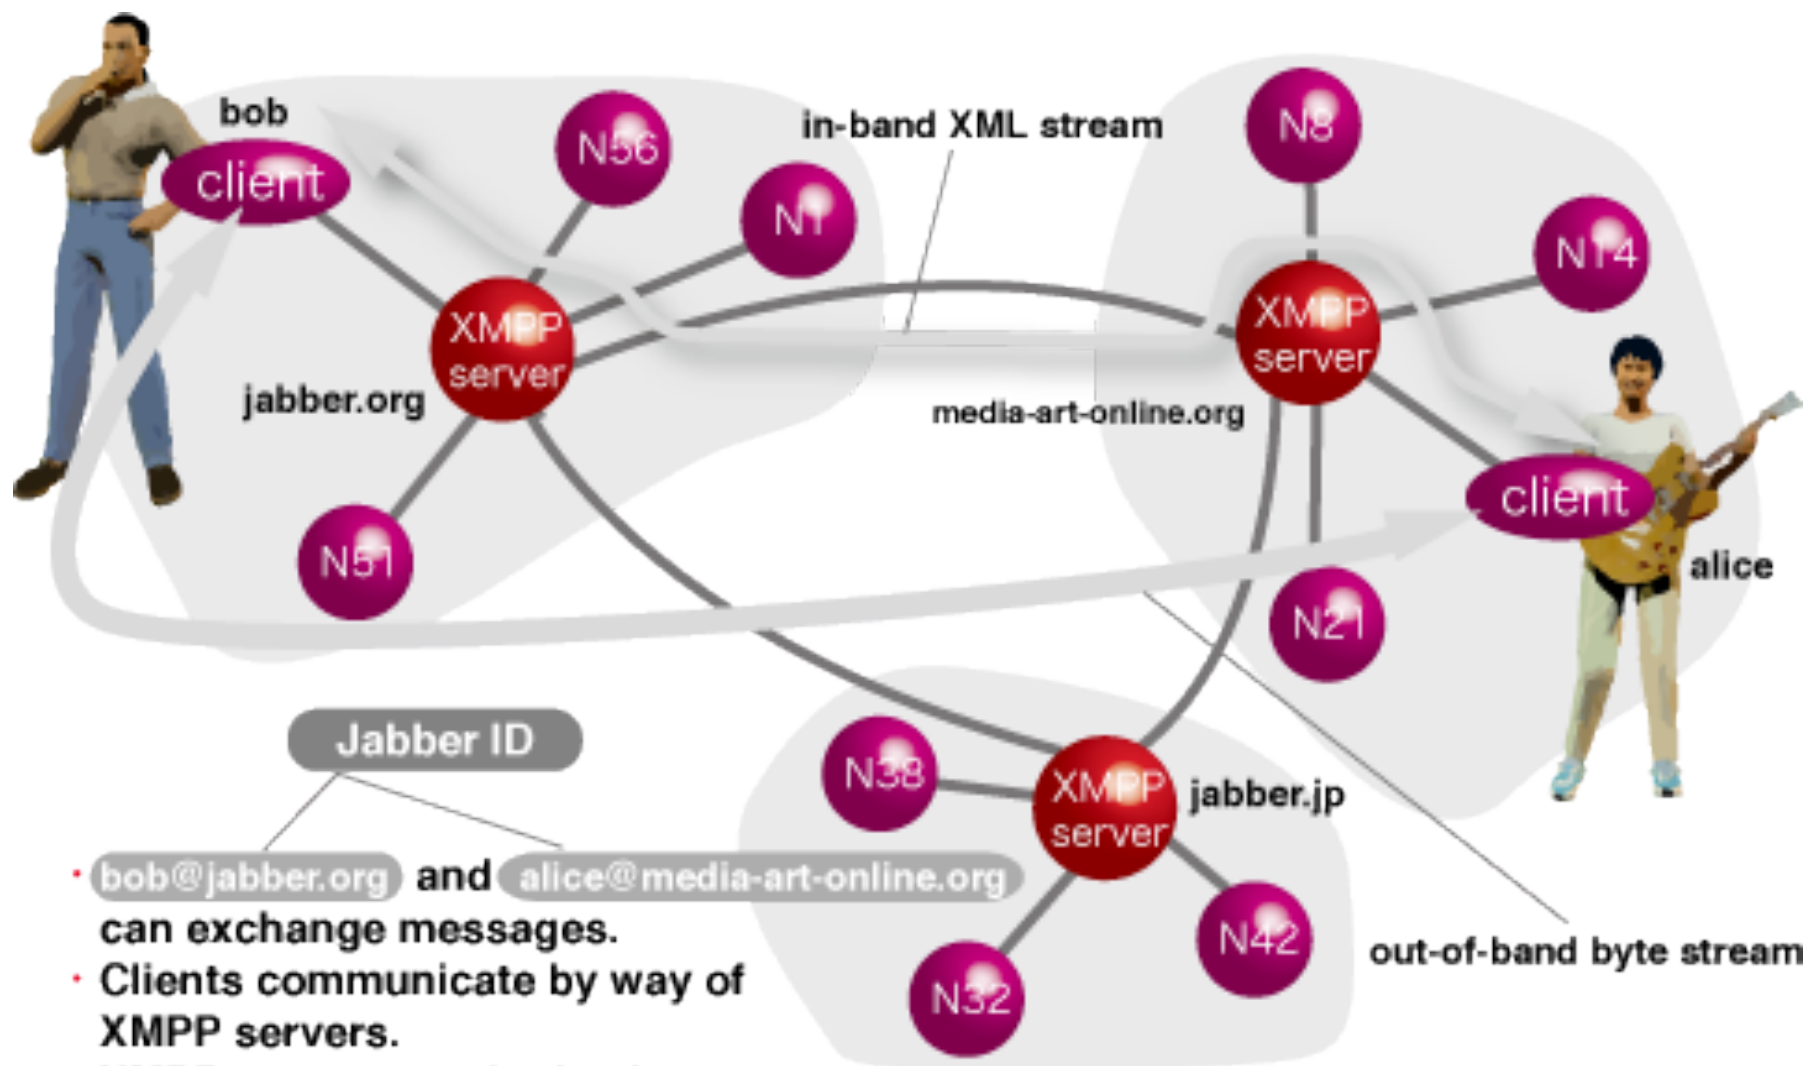
- Higher layers
- HTTP
- XMPP
 - Actors
 - Protocol
 - Extensions: BOSH
- Radically different

XMPP

- EXtensible Messaging & Presence Protocol
 - Basis for Jabber and Google Talk
 - RFC 3920, 3921
- Core provides XML streaming
 - Presence, Messages, Iq
- Many extensions
 - HTTP binding, real-time media signalling, multi-user chat, publish-subscribe, &c.
 - XEP series

Actors

- *Clients*, connect to ...
- *Servers*, which may interconnect directly
 - Servers are networked and can route messages
 - Typically multiplexed over single TCP connection
- *Gateways*, connect foreign clients to XMPP network
 - E.g., gateway Skype into a Google Talk session
- Address format is JID: node@domain/resource
 - Bare JID drops the /resource
- Security via TLS



- **bob@jabber.org** and **alice@media-art-online.org** can exchange messages.
- Clients communicate by way of XMPP servers.
- XMPP servers can be freely set up.

Protocol

- Exchange *XML streams*
 - Multiple *XML stanzas*
 - ...encapsulated in stream
- Three stanzas defined
 - presence
 - Express availability
 - Negotiate subscriptions
 - message
 - Push information
 - iq
 - Request-response

```
<stream>
```

```
<presence>
```

```
<show/>
```

```
</presence>
```

```
<message to='foo'>
```

```
<body/>
```

```
</message>
```

```
<iq to='bar'>
```

```
<query/>
```

```
</iq>
```

```
...
```

```
</stream>
```

Example

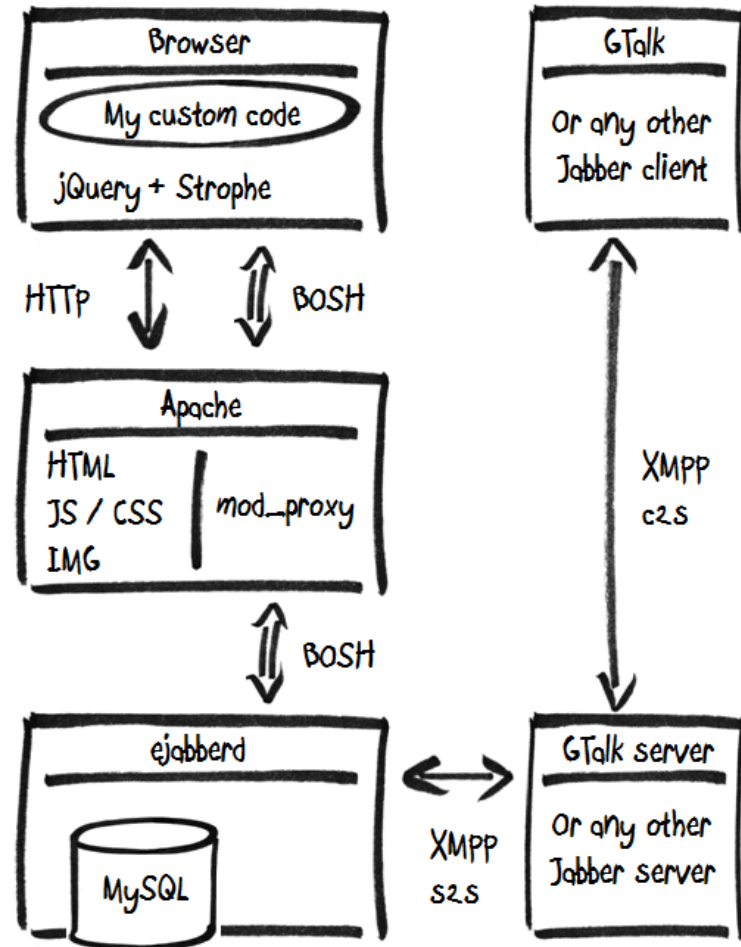
```
C: <?xml version='1.0'?>
  <stream:stream
    to='example.com'
    xmlns='jabber:client'
    xmlns:stream='http://etherx.jabber.org/streams'
    version='1.0'>
S: <?xml version='1.0'?>
  <stream:stream
    from='example.com'
    id='someid'
    xmlns='jabber:client'
    xmlns:stream='http://etherx.jabber.org/streams'
    version='1.0'>
... encryption, authentication, and resource binding ...
C: <message from='juliet@example.com'
  to='romeo@example.net'
  xml:lang='en'>
C:   <body>Art thou not Romeo, and a Montague?</body>
C: </message>
S: <message from='romeo@example.net'
  to='juliet@example.com'
  xml:lang='en'>
S:   <body>Neither, fair saint, if either thee dislike.</body>
S: </message>
C: </stream:stream>
S: </stream:stream>
```


Extensions

- Administered by the XMPP Foundation
 - <http://xmpp.org/extensions/>
 - E.g., XEP-0124, BOSH; XEP-0206, XMPP over BOSH
- Bidirectional Streams Over Synchronous HTTP
 - Firewall friendly (TCP/80 vs. TCP/{5222-3, 5269})
 - Free compression (most servers support Gzip)
 - Hides unreliability
 - Emulates long-lived connection by a sequence of request-responses

BOSH

- Naïve:
Client periodically polls server
 - High latency and wastes bandwidth and battery
 - Matters especially on mobile clients
- Better:
Send new request on receipt
 - Server always has outstanding connection down which it can push data
- Works well with HTTP/1.1 but not so bad with 1.0



Corner Cases

- Client gets new data
 - Existing connection is blocked at the server
 - So open new connection to send, causing server to close old one
- Nothing happens for several minutes
 - Need a *keepalive*
 - Server returns empty response and client sends a new empty request
- Constrained client
 - Can't do HTTP/1.1 or multiple connections
 - Revert to naive polling mode

Contents

- Higher layers
- HTTP
- XMPP
- Radically different
 - Peer-to-peer
 - BitTorrent
 - Active networks

Peer-to-Peer (P2P)

- Both previous protocols were client-server
 - What if the server is the bottleneck?
 - ...whether through CPU, network, management...
- Alternative: peer-to-peer systems
 - E.g., CAN, CHORD, Pastry, BitTorrent, KaZaA, &c.
 - No designated central point (but consider BT *tracker*)
 - Typically self-organizing
- Often provide *distributed hash table* abstraction
 - Structured vs. unstructured
 - Message delivery usually scales as $\log(N)$ hops with network size N

BitTorrent

- Distributes load away from a single source site
 - (Approx.) proportional to popularity
- File divided into *pieces*, obtained separately
 - In random order, trying to keep *file* live
 - Each piece has a hash to provide integrity
- *Torrent descriptor* file made available to a *seed*
- *Tracker* knows who's participating in torrent
 - Can itself be distributed, e.g., DHT methods
- Client contacts tracker to obtain list of peers
 - Connect to peers, start downloading pieces
 - “Fair” to use many TCP connections to download?

Common P2P Issues

- Seeding the swarm: problems of flash crowds
- Anger of netadmins: breaks usage models
- No anonymity: leaves you open to attack
- Validity of metadata: bad torrents
- Leeching
 - Tit-for-tat schemes penalise newcomers
 - ...but what are the incentives for peers to share?
- The Law: but is a protocol illegal?

Active Networks

- And now for something completely different
 - Depending on how you look at it
- The network considers packets to be passive
 - Routers and middleboxes just forward
 - ...with a bit of rewriting, possibly triggering response
- What about if each packet was a bit of code?
 - Routers execute packet (header?) instead
 - Interesting research idea, never really took off
 - Lots of cool stuff about constraining runtime environment, proving properties on the code, &c.

Summary

- In the IP stack, session and presentation layers are generally subsumed into application layer
- Two widely used and interesting application protocols are HTTP and XMPP
- Boundless possibility, e.g., peer-to-peer active networks

Quiz

1. What do the session and presentation layers provide?
2. What characterises a P2P protocol?
3. XMPP allows two clients to exchange messages. Why is it **not** described as a P2P protocol?
4. How do proxies complicate HTTP? Why are they still used?
5. Describe the benefits of pipelining in HTTP/1.1 with an example.
6. What security guarantees does TLS provide, and how?
7. Why is BOSH required to run XMPP over HTTP?