

Programming

G54ACC

Lecture 16

richard.mortier@nottingham.ac.uk

Recap

- Basics of Berkeley Sockets API
 - `socket()`, `connect()`, `send()`, `recv()`, &c
- Endianness
- Multiplexing at lower layers
 - We'll discuss multiplexing from the programmer point of view
- Remote Procedure Call (RPC)
 - Invocation of code on a remote machine

Contents

- Berkeley Sockets
- Web RPC
- RESTful web services
- Flavours of “cloud computing”
- Summary

Contents

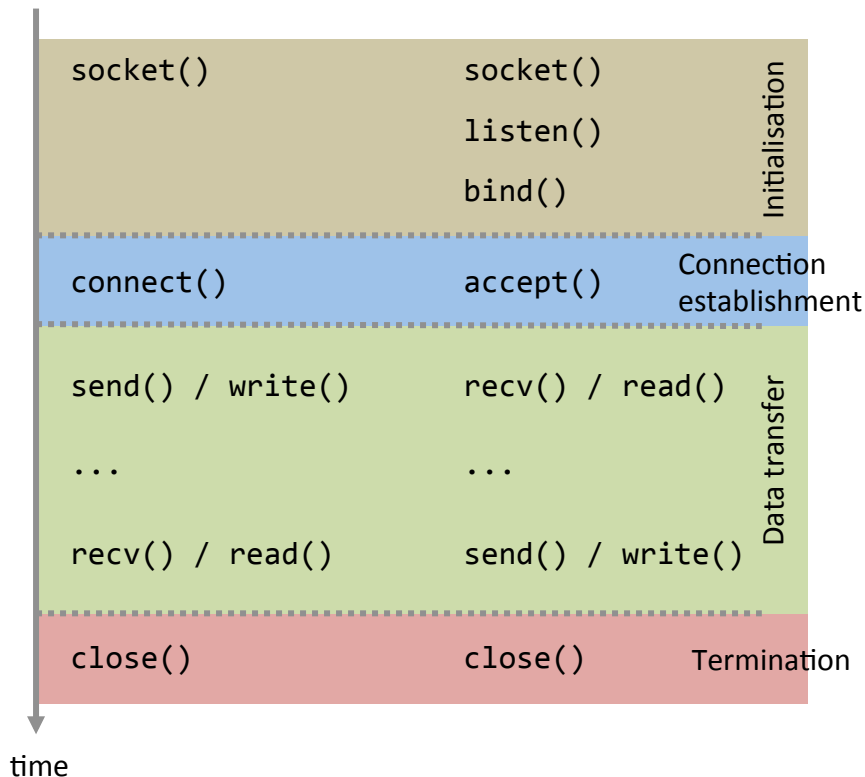
- Berkeley Sockets
 - TCP/UDP
 - Endianness
 - Multiplexing
- Web RPC
- RESTful web services
- Flavours of “cloud computing”
- Summary

Sockets

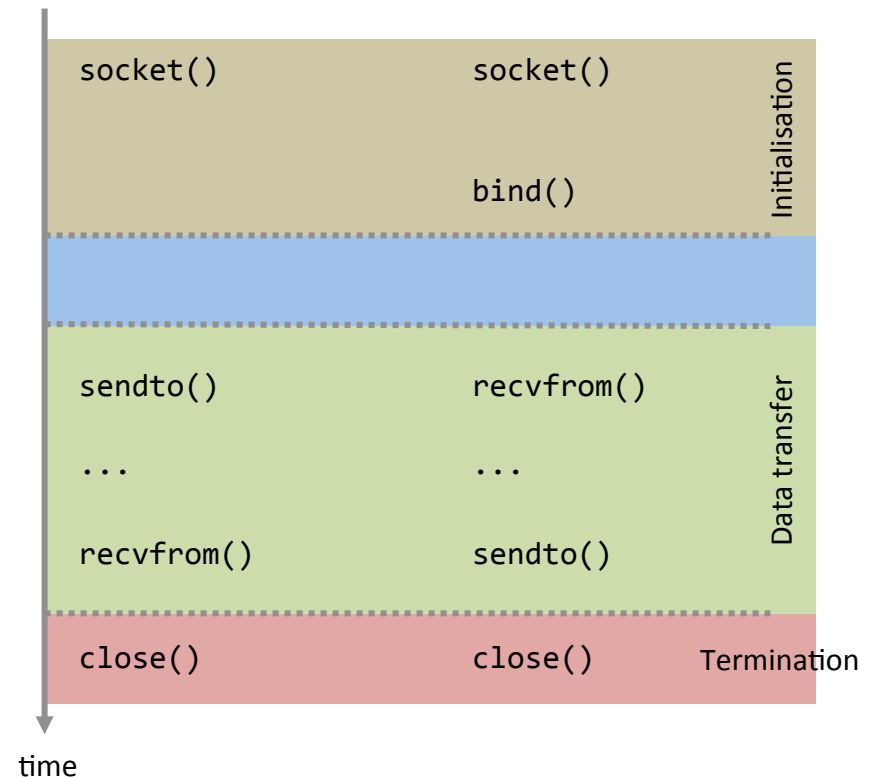
- Berkeley Sockets API, 4.2BSD (1983)
 - Use *Internet sockets*, presented as file descriptors
 - Extended to many other protocols
- `int socket(int domain, int type, int protocol)`
 - domain: communication domain
 - PF_INET, PF_INET6, PF_LOCAL, PF_RAW
 - type: communication semantics
 - SOCK_STREAM, SOCK_DGRAM, SOCK_RAW
 - protocol: usu. 1:1 with (domain,type)
 - IPPROTO_ICMP, IPPROTO_TCP, IPPROTO_UDP

TCP vs. UDP

- File-like interface
 - Streaming



- Packet-like interface
 - Need to packetize



Endianness


- Host byte order need not be identical
 - Big endian (most-significant-byte first)
 - vs. Little endian (least-significant-byte first)
- Need to decide which we use on the network
 - Lest 0x0001 be received as 0x0100
 - Or negotiate every time?!
- Network byte order == big endian
 - No reason. It just is.
 - Convert via `u[16|32] htonl[s|1]/ntoh[s|1](u[16|32])`

<http://flickeringtubelight.net/blog/2004/05/big-endian-and-little-endian-storage-schemes-how-to-remember/>

base 10: 3735928559 = base 16: 0xDEADBEEF


In Big Endian storage schemes, the most significant byte is stored "first" - i.e. "big end" goes to the lower address. Memory is addressed/accessed from low to high addresses

base 2: 11011110 10101101 10111110 11101111



In Little Endian storage schemes, the least significant byte is stored "first" - i.e. "little end" goes to the lower address. Memory is addressed/accessed from low to high.

base 2: 11011110 10101101 10111110 11101111



Multiplexing

- Naive: read and block
 - But what if you need to handle >1 socket?
- Naive 2: poll non-blocking socket
 - Wastes CPU
- Use `select()`, either blocking or non-blocking
 - Enables waiting (until timeout) on any socket of a set becoming active
 - `int select(int nfds, fd_set *rfds, fd_set *wfds, fd_set *efds, struct timeval timeout)`
- Check W. Richard Stevens “*Unix Network Programming*” and “*TCP/IP vol.1*” for details

Contents

- Berkeley Sockets
- **Web RPC**
 - XML-RPC
 - SOAP
- RESTful web services
- Flavours of “cloud computing”
- Summary

Web RPC Systems

- (At least) Two: XML-RPC *vs.* SOAP
 - Appears to be/have been a religious war here
 - SOAP evolved from XML-RPC
- Both provide
 - Data encoded as XML
 - Remote Procedure Call interfaces
 - ...to remote web services

XML-RPC

- Older, simpler
 - <http://www.xmlrpc.com/spec>
- Community based

```
POST /RPC2 HTTP/1.0
User-Agent: Frontier/5.1.2 (WinNT)
Host: betty.userland.com
Content-Type: text/xml
Content-length: 181
```

```
<?xml version="1.0"?>
<methodCall>
  <methodName>examples.getBounds</methodName>
  <params>
    <param>
      <value><i4>41</i4></value>
    </param>
  </params>
</methodCall>
```

```
HTTP/1.1 200 OK
Connection: close
Content-Length: 158
Content-Type: text/xml
Date: Fri, 17 Jul 1998 19:55:08 GMT
Server: UserLand Frontier/5.1.2-WinNT
```

```
<?xml version="1.0"?>
<methodResponse>
  <params>
    <param>
      <struct>
        <member>
          <name>lowerBound</name>
          <value><i4>18</i4></value>
        </member>
        <member>
          <name>upperBound</name>
          <value><i4>139</i4></value>
        </member>
      </struct>
    </param>
  </params>
</methodResponse>
```

SOAP

- Newer, more featureful, more flexible
 - Doesn't require HTTP as transport
 - Provides e.g., discovery (WSDL), &c
- Not particularly simple
- W3C standard

```
POST /InStock HTTP/1.1
Host: www.example.org
Content-Type: application/soap+xml;
charset=utf-8
Content-Length: nnn

<?xml version="1.0"?>
<soap:Envelope
xmlns:soap="http://www.w3.org/2001/12/
soap-envelope"
soap:encodingStyle="http://www.w3.org/
2001/12/soap-encoding">

  <soap:Body xmlns:m="http://
www.example.org/stock">
    <m:GetStockPrice>
      <m:StockName>IBM</m:StockName>
    </m:GetStockPrice>
  </soap:Body>

</soap:Envelope>
```

Contents

- Berkeley Sockets
- Web RPC
- RESTful web services
 - REST constraints and interfaces
 - Creating a RESTful service
- Flavours of “cloud computing”
- Summary

RESTful Web Service

- Roy Fielding during HTTP/1.1 development (1994–2002)
 - doi:10.1145/514183.514185
 - http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm
- Goals
 - Minimize latency, communication
 - Maximise implementation scalability, component independence

RESTful Web Service

- A RESTful service conforms to the REST constraints
 - Representational State Transfer
 - Makes more extensive use of HTTP features
 - In principle can use other transports
- Client requests generate server responses
- Both built around *representations of resources*
 - HTML, JPEG, whatever in principle. JSON or XML in practice.

REST Constraints

- Six architectural constraints
 - Client-server, Cacheable, Layered
 - Separate concerns, beware existence of caches, proxies, &c
 - Code on demand (optional)
 - Server can push e.g., JavaScript to client as part of response
 - Stateless, Uniform interface
 - ...

Stateless

- Does *not* mean the server must be stateless!
- Simply that no *client* context is stored on the server between requests
 - Each client request contains enough to service it
 - If server does store state, it is addressable as a resource via URL in the usual way
- Contrast

GET /resources/nextPage?

...

```
<?xml version="1.0"?>
<rsp status="ok">
  <resource id="1" />
  <resource id="2" />
</rsp>
```

GET /resources/?page=2

...

```
<?xml version="1.0"?>
<rsp page="2" nextPage="3">
  <resource id="21" />
  <resource id="22" />
</rsp>
```

Uniform Interface

- Client uses representation to modify server state, addressed via URL
- RESTful service maps CRUD to HTTP methods
 - CREATE: POST
 - READ/RETRIEVE: GET, idempotent, side-effect free
 - UPDATE: PUT, idempotent (since uses specified name)
 - DELETE: idempotent
- CREATE/UPDATE as POST vs. PUT appears debated
 - Differ in treatment of Request-URI and request entity
 - POST: treat entity as “new subordinate”
 - PUT: entity “stored under supplied Request-URI”

Creating a RESTful Service

- What are the URIs?
 - Find the nouns, define *collections*
- What's the format?
 - XML, JSON, &c
- What methods are supported at each URI?
 - Mapping to the CRUD functions
- What status codes could be returned?
 - 1xx (informational), 2xx (successful), 3xx (redirection), 4xx (client error), 5xx (server error)

Notable Points

- Collections
 - <http://.../people> : GET (list), POST (append), PUT (replace), DELETE
 - <http://.../people/1> : GET (retrieve), PUT (update), DELETE
- Tickets used to provide *at-most-once* semantics
 - Unique, server-generated GUIDs
 - POST to /ticket/o to create a “slot” in the server
- Redirects help clients handle tickets
 - Ticket response redirects to /o/<n>
 - PUT /o/<n> creates new item, idempotently

Contents

- Berkeley Sockets
- Web RPC
- RESTful web services
- Flavours of “cloud computing”
- Summary

Cloud Computing

- Datacenter as a utility
 - Buy it (rent it) when you need it (OpEx vs. CapEx)
 - Scales “on demand”
- Beware!
 - Variable performance
 - Little control over data placement
 - Unknown trust implications
 - Cloud provider as single-point-of-failure
 - Not just their hosts, but their and your network connectivity
 - Building scalable services is still hard!

Flavours

- X as a Service for X =
 - Software: Salesforce
 - Provides you with a Customer Relationship Manager service
 - No longer need to manage software, upgrades, &c
 - Platform: Google App Engine, PiCloud, Heroku, Azure
 - Hosts web apps backed by shared store (SQL-like)
 - Fires up (Python, Java, Ruby, CLR) VM as required
 - Infrastructure: Amazon EC2, Rackspace
 - Provides XEN hosted VMs (Linux, Windows, custom)
 - Also persistent store (S3), other infrastructure APIs

Summary

- There are a wide range of APIs that attempt hide the complexity of network programming
 - Sockets is the (old) canonical Internet API on which all else is built
 - Web RPC was briefly popular
 - RESTful platforms with APIs are “web2.0” and all the rage (and jolly useful)
- Cloud computing in some ways goes a stage further, hiding even more of the complexity
 - But you introduce other problems due to, e.g., the scale of the application
- Failures are always a pain to deal with, anywhere