

Virtualization

G54CCS – Lecture 4

Richard Mortier

<http://www.cs.nott.ac.uk/~rmm/teaching/2011-g54ccs/>

Don't forget!

- Extra lab session **today** 4pm, C11
- Appears that B52 also free at that time
- Please split yourselves between the two

Overview

- Why Virtualization?
- Abstraction & Layering
- The Operating System
- Virtualization
- Benefits

Overview

- Why Virtualization?
- Abstraction & Layering
- The Operating System
- Virtualization
- Benefits

Why Virtualization?

- Perhaps the key enabling cloud technology
 - How?
 - Why?
- To answer, we must dive into the technology
 - (a little)
- Three important concepts:
 - Abstraction/Layering
 - Isolation
 - Multiplexing

Overview


- Why Virtualization?
- Abstraction & Layering
 - What's Under Your Application?
 - Hardware
- The Operating System
- Virtualization
- Benefits

Abstraction and Layering

- Computers are complex – need to control this
- *Abstraction*
 - Deliberately hiding details
 - Presenting a simplified *interface*
- *Layering*
 - Control interactions
 - Each layer depends only on interface below
 - ...and services layer above via own interface


What's Under Your App?

- Application code
 - E.g., in your labs
 - What you want to run
 - What handles your users' requests
 - What makes you money




What's Under Your App?

- Application code
 - E.g., in your labs
 - What you want to run
 - What handles your users' requests
 - What makes you money
- Sits on top of a *runtime*
 - python.exe; JVM; libc
 - Provides environment and support functions



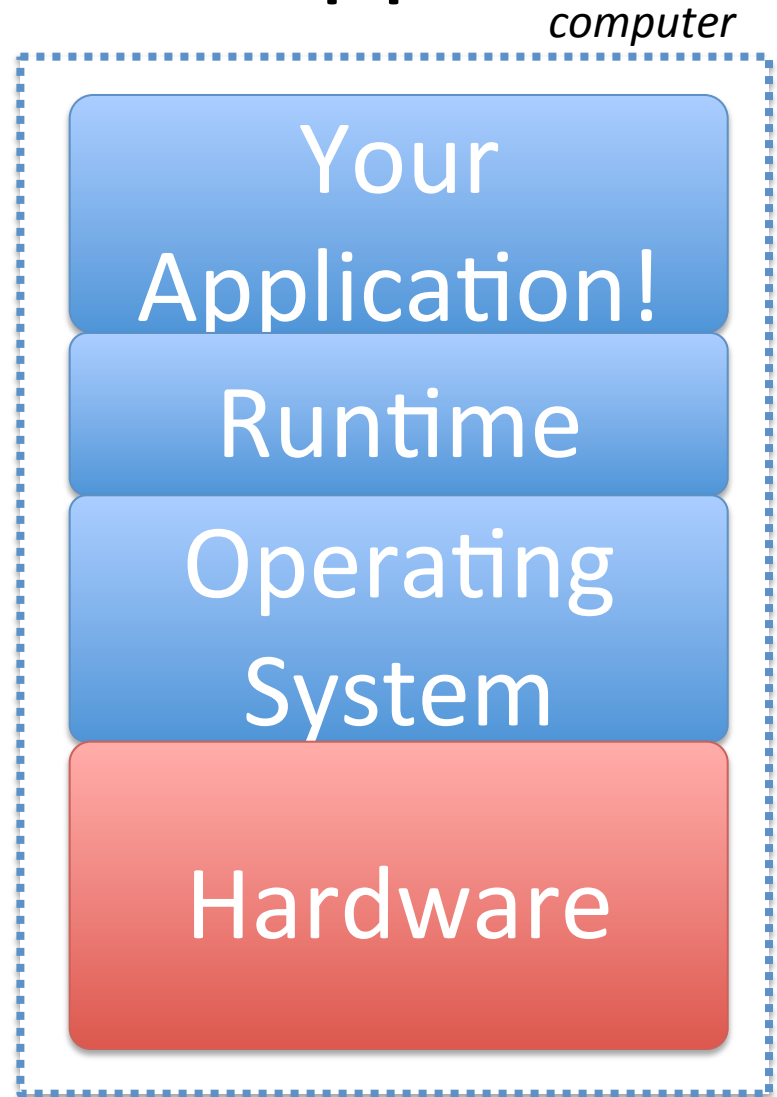
What's Under Your App?

- Operating System
 - Abstracts hardware
 - Isolates code
 - Manages resources



What's Under Your App?

- Operating System
 - Abstracts hardware
 - Isolates code
 - Manages resources
- Hardware
 - CPU with N cores
 - Memory, storage
 - Communications



Hardware

- The physical components of the computer
- Processing
 - CPU, motherboard, GPU
- Memory
 - Registers, cache, main memory
- Peripherals
 - Communications – network cards
 - Storage – Hard disks, SSD
 - Screens, cameras, printers, ...

Overview

- Why Virtualization?
- Abstraction & Layering
- The Operating System
 - The Boot Process
 - Resource Management
 - Isolation, Buffering, Multiplexing
- Virtualization
- Benefits

The Boot Process (simplified!)

- Power on and stable
- CPU starts executing code
- Goes first to the BIOS
 - Initiates POST
 - Finds boot device, loads MBR code, jumps
- MBR code knows how to read disk
 - Find, load, execute operating system
- Can now run application code
 - E.g., a shell, a window manager

The Operating System (simplified!)

- Minimally,
 - Manages resources
 - Simplifies hardware
 - Provides application programming interface (API)
- As part of *resource management*,
 - Isolation
 - Buffering
 - Multiplexing

Resources

- CPU
 - What should run now?
- Memory
 - Hide memory hierarchy from programmer
 - Virtual address space
- Storage
 - OS provides high level interface to block storage
 - E.g., file system, key-value store, object serialization
- Network
 - Range of communications protocols and primitives
 - Access to hardware to send and receive traffic

Isolation

- Hardware isolation – one machine, one app
 - But typically even then, many pieces of software
 - E.g., OS, standard services, application
- Software isolation
 - Use *processes* to protect memory between apps
 - Other concurrency models are available
 - Limits scope for interference and corruption
 - Application may have many processes

Interacting with Hardware


- How to protect interactions with hardware?
 - Prevent untrusted user code having direct access
- Kernel vs. User mode
 - Use hardware support – x86 *rings*
 - *syscalls* provide interface to OS
 - E.g., `read()`, `write()`, `recv()`, `send()`, ... [POSIX]
- Manage passing of buffers across boundary
 - Must ensure that app can't corrupt memory

Multiplexing

- Multiple applications running simultaneously
 - ...or the illusion of that if only 1 CPU
- Makes more efficient use of the hardware
 - While one app is waiting, another is getting on
 - If none are waiting, then switch between apps
 - Wide range of policies for who goes next
- Waiting is quite common due to latency mismatch
- Example – Web server waiting on database
 - Request comes in, gets processed
 - Response needs several DB records plus some files


Capacity vs. Latency

approximate!



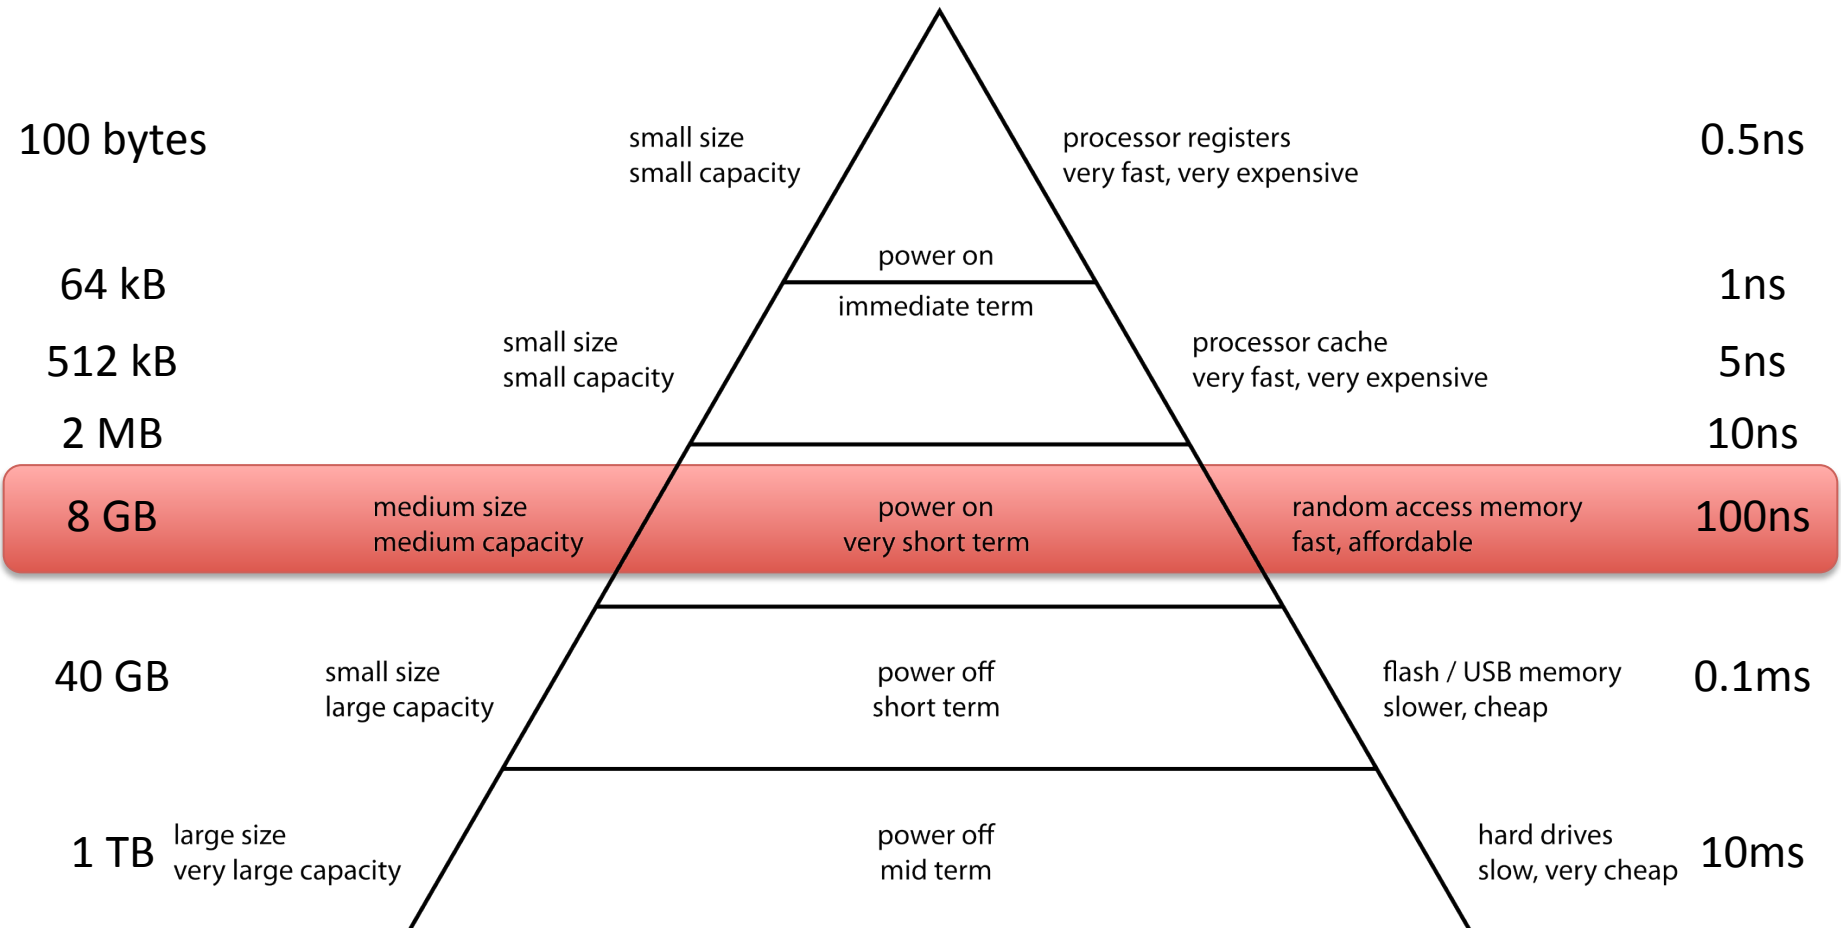
Capacity vs. Latency

approximate!




Capacity vs. Latency

approximate!




Capacity vs. Latency

approximate!



Capacity vs. Latency

approximate!



Overview

- Why Virtualization?
- Abstraction & Layering
- The Operating System
- Virtualization
 - Virtualization Types
 - Compatibility Options
 - In the Wild
- Benefits

Virtualization

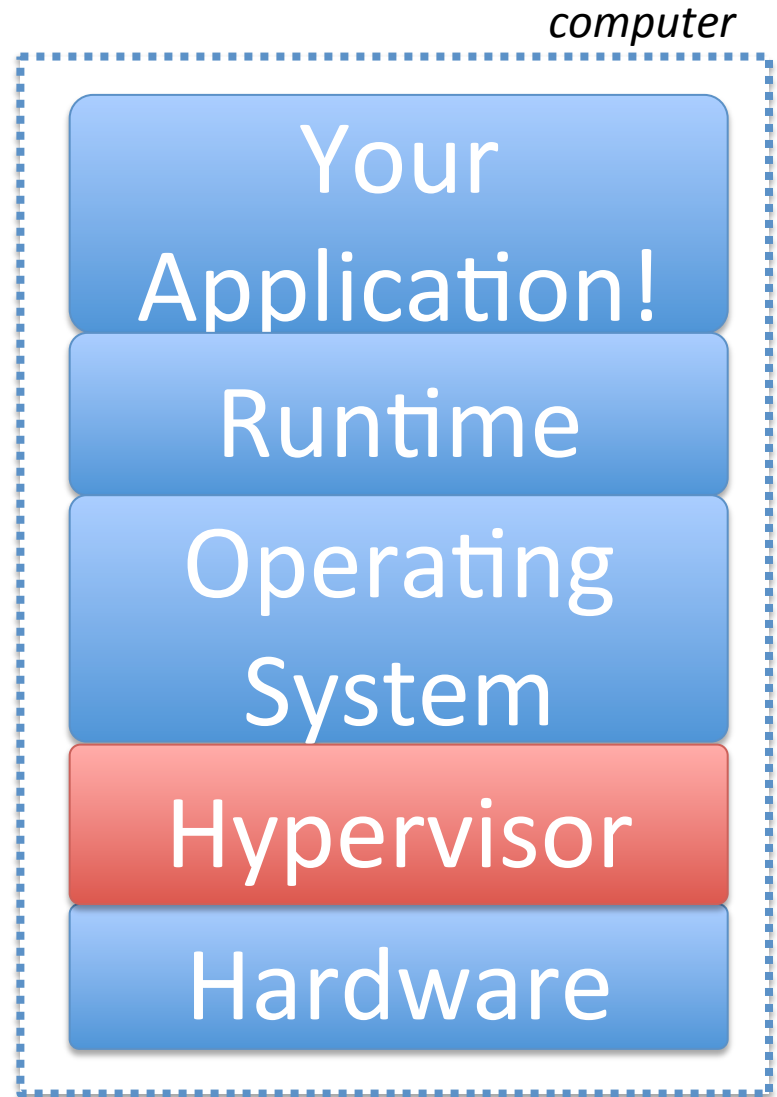
- Software deployment
 - More than just the application
 - Applications rely on many services
 - E.g., DNS, firewall, filtering, backup, storage
- Can we do better?

Virtualization

- Software deployment
 - More than just the application
 - Applications rely on many services
 - E.g., DNS, firewall, filtering, backup, storage
- Can we do better?
 - Encapsulate OS and associated state
 - Configuration, data, filesystem, code
- But how?

Virtualization

- Apply *indirection*
 - A traditional computer science solution
 - Insert extra layer of code to support the OS
 - OS itself still supports application processes
 - *Virtual Machine (VM)*



Virtualization Types

- Type I – *Hypervisor Architecture*
 - So-called “bare metal”
 - A *hypervisor* or *VMM* sits beneath the OS
 - On x86, OS must be modified to target this
 - E.g., Xen, VMWare ESX Server
- Type II – *Hosted Architecture*
 - Sits above a *host OS*
 - E.g., VirtualBox, VMWare Workstation

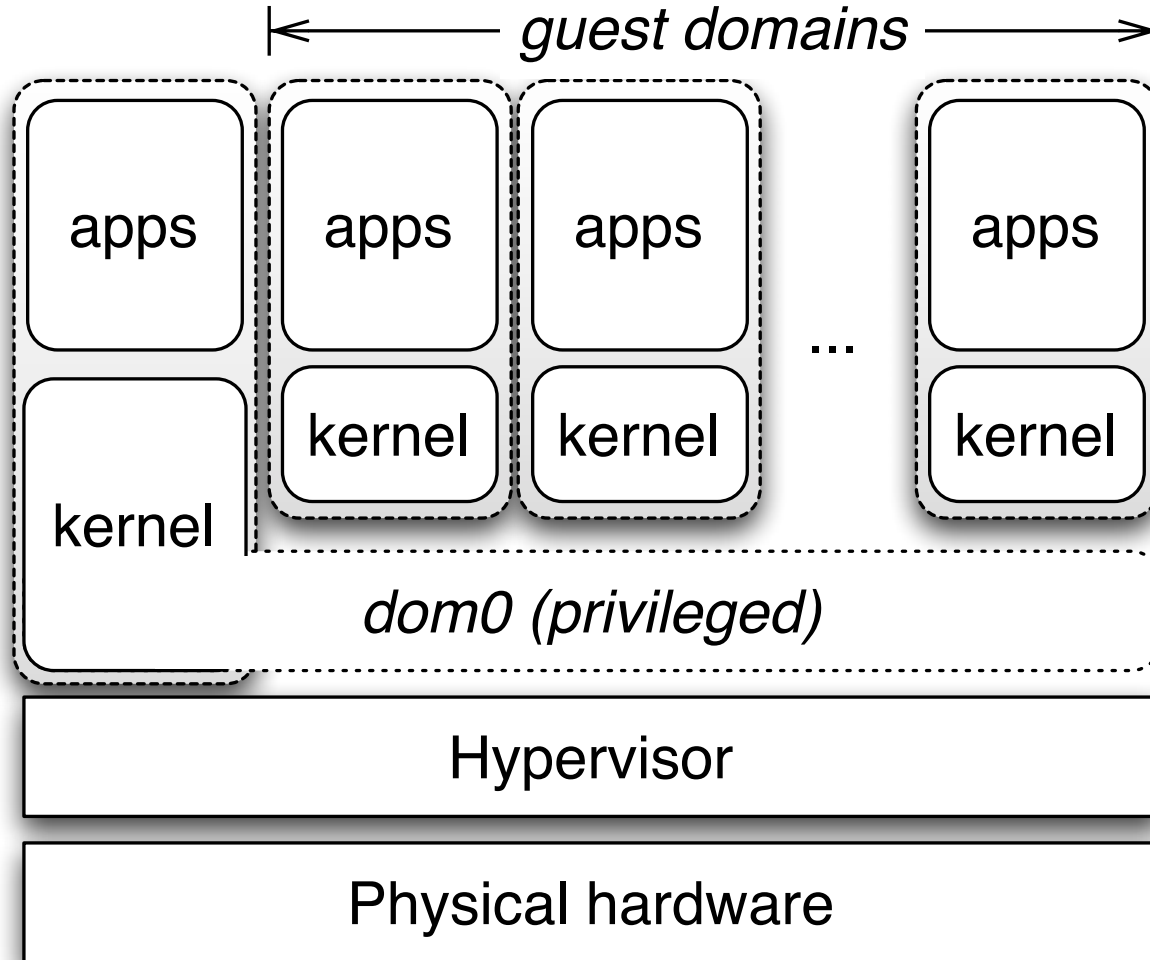
Compatibility Options

- Full virtualization
 - Completely transparent to guest OS
 - Binary level compatibility, no changes required
 - Potential performance issues
- Paravirtualization
 - Transparent to applications
 - Requires modification of guest OS
 - Xen required ~3k lines of Linux
 - Improves performance

Aside – x86 Virtualization

- Instruction set difficult to virtualize
 - Non-privileged instructions expose VMM
 - Three approaches
- Emulation (VirtualPC)
 - Emulate all instructions in kernel mode
- Binary translation (VMWare)
 - Convert kernel code to call into VMM
- Paravirtualization (Xen)
 - Remove privileged instructions from Guest OS

Example – Xen



In the Wild

- Amazon's EC2 infrastructure
 - Uses a (patched) Xen hypervisor
 - Supports Linux, Windows, other VM images
 - Amazon Machine Image (AMI) as self-contained VM
 - Can provide own kernels too nowadays
- Microsoft's Azure platform
 - Customized Hyper-V
 - Provides wide range of Windows VM images, plus selected Linux distributions
 - Had to submit code for drivers for Linux

Overview

- Why Virtualization?
- Abstraction & Layering
- The Operating System
- Virtualization
- **Benefits**

Benefits

- Natural provisioning primitive
 - Easier to deploy and manage
 - Take advantage of features like live migration
- Datacenter scaledown
 - Coalesce machine roles onto less hardware
 - Reduces costs in all directions
- Security
 - Better isolation can improve security, resilience
- Charging
 - Separate resource consumption, separate billing

Summary

- Why Virtualization?
- Abstraction & Layering
 - What's Under Your Application?
 - Hardware
- The Operating System
 - The Boot Process
 - Resource Management
 - Isolation, Buffering, Multiplexing
- Virtualization
 - Virtualization Types
 - Compatibility Options
 - In the Wild
- Benefits